



PRESIDENCY UNIVERSITY

Private University Estd. in Karnataka State by Act No. 41 of 2013

Itgalpura, Rajankunte, Yelahanka, Bengaluru – 560064



RURAL-MED NEXUS: AN OFFLINE INTELLIGENT MEDICAL KNOWLEDGE SYSTEM

A PROJECT REPORT

Submitted by,

MOHAMED ASLAM PASHA – 20221ISE0053

KOVURI NIZAMUDDIN – 20221ISE0083

VASEEM B M – 20221ISE0084

Under the guidance of,

Mr. SANTHOSH KUMAR K L

BACHELOR OF TECHNOLOGY IN INFORMATION SCIENCE AND ENGINEERING (AI AND ROBOTICS)



GAIN MORE KNOWLEDGE
REACH GREATER HEIGHTS

PRESIDENCY UNIVERSITY

BENGALURU, APRIL 2026



PRESIDENCY UNIVERSITY

Private University Estd. in Karnataka State by Act No. 41 of 2013

Itgalpura, Rajankunte, Yelahanka, Bengaluru – 560064



PRESIDENCY SCHOOL OF COMPUTER SCIENCE AND ENGINEERING

BONAFIDE CERTIFICATE

Certified that this report “Rural-Med Nexus: An Offline Intelligent Medical Knowledge System” is a Bonafide work of “Mohamed Aslam Pasha (20221ISE0053), Kovuri Nizamuddin (20221ISE0083), Vaseem B M (20221ISE0084)”, who have successfully carried out the project work and submitted the report for partial fulfilment of the requirements for the award of the degree of BACHELOR OF TECHNOLOGY in INFORMATION SCIENCE AND ENGINEERING (AI AND ROBOTICS) during 2025-26.

Mr. Santhosh Kumar K L

Assistant Professor

Project Guide

PSCS

Presidency University

Ms. Sumita Guddin

Program Project Coordinator

PSCS

Presidency University

Dr. Gopal Krishna Shyam

Head of the Department,

PSCS-COM, ISE.

Presidency University

Dr. Duraipandian N

Dean

PSCS

Presidency University

Examiners

Sl. no.	Name	Signature	Date
1	Mr. Santhosh Kumar K L		
2	Mr. Vivek Bongale		

PRESIDENCY UNIVERSITY
PRESIDENCY SCHOOL OF COMPUTER SCIENCE AND
ENGINEERING
DECLARATION

We the students of final year B.Tech in INFORMATION SCIENCE AND ENGINEERING (AI AND ROBOTICS) at Presidency University, Bengaluru, named Mohamed Aslam Pasha, Kovuri Nizamuddin, Vaseem B M, hereby declare that the project work titled “Rural-Med Nexus: An Offline Intelligent Medical Knowledge System” has been independently carried out by us and submitted in partial fulfilment for the award of the degree of B-Tech in INFORMATION SCIENCE ENGINEERING (AI AND ROBOTICS), during the academic year of 2025-26. Further, the matter embodied in the project has not been submitted previously by anybody for the award of any Degree or Diploma to any other institution.

Mohamed Aslam Pasha	USN: 2022IISE0053
Kovuri Nizamuddin	USN: 2022IISE0083
Vaseem B M	USN: 2022IISE0084

PLACE: BENGALURU

DATE:



PRESIDENCY UNIVERSITY

Private University Estd. in Karnataka State by Act No. 41 of 2013
Itgalpura, Rajankunte, Yelahanka, Bengaluru – 560064



ACKNOWLEDGEMENT

For completing this project work, I have received support and guidance from many individuals, to whom I would like to express my deep sense of gratitude and indebtedness. I extend my sincere gratitude to our beloved **Chancellor, Vice Chancellor, Pro-Vice Chancellor, and Registrar** for their support and encouragement in the completion of the project.

I express my cordial thanks to **Dr. Duraipandian N, Dean, Presidency School of Computer Science and Engineering**, and the Management of Presidency University for providing the required facilities and an intellectually stimulating environment that aided in the completion of my internship work.

I am also thankful to **Dr. Gopal Krishna Shyam, Professor and Head of the Department**, Presidency School of Computer Science and Engineering, Presidency University, for his mentorship and encouragement.

I would like to sincerely thank my internal guide **Mr. Santhosh Kumar K L, Assistant Professor**, Presidency School of Computer Science and Engineering, Presidency University, for the moral support, motivation, timely guidance, and encouragement provided during the period of project work.

I am grateful to **Ms. Sumita Guddin, Program Project Coordinator**, Presidency School of Computer Science and Engineering, for facilitating problem statements, coordinating reviews, monitoring progress, and providing valuable support and guidance.

I also extend my sincere thanks to the **Teaching and Non-Teaching staff** of the Presidency School of Computer Science and Engineering, and staff from other departments, for their valuable help and cooperation.

MOHAMED ASLAM PASHA

KOVURI NIZAMUDDIN

VASEEM B M

ABSTRACT

Limited medical infrastructure, a shortage of trained professionals, and lack of reliable internet access continue to restrict healthcare delivery in rural and connectivity-limited regions. Modern AI-powered medical assistants rely on cloud APIs, which leads to privacy risks and limits functionality in remote environments. In this project we present Rural-Med Nexus, a fully offline, privacy-preserving multimodal RAG system tailored towards rural healthcare ecosystems.

It then allows for ingestion of medical data from PDFs, DOCX, TXT, images (using OCR), and audio (using ASR), constructs chunks of meaning semantically producing embeddings using transformer based models before indexing them in FAISS vector search for rapid similarity retrieval. To minimize hallucination risks, a lightweight local LLM generates source-grounded medical responses by using retrieved context. In addition, the architecture is completely internet-free and designed for performance with a CPU in low-resource settings.

Rural-Med Nexus guarantees that patient data stays solely within the local device, facilitates citing and referencing based on metadata, and generates medically responsible responses tailored to rural healthcare contexts. This system combines edge AI with retrieval grounding and helps toward secure, scalable and equitable access to medical intelligence for these underserved communities.

TABLE OF CONTENTS

Sl. No.	Title	Page No.
	Declaration	III
	Acknowledgement	IV
	Abstract	V
	List of Figures	IX
	List of Tables	XI
	Abbreviations	XII
1.	Introduction	1
	1.1 Background	2
	1.2 Statistics	3
	1.3 Prior existing technologies	4
	1.4 Proposed approach	5
	1.5 Objectives	6
	1.6 SDGs	8
	1.7 Overview of project report	9
2.	Literature review	12
3.	Methodology	27
	3.1 Ingestion Pipeline	27
	3.2 Query Pipeline	29
	3.3 Retrieval-Augmented Generation (RAG) Strategy	30
	3.4 V-Model	31
	3.5 Key Technologies and Tools Used	32
	3.6 Benefits of the Proposed System	33

	3.7 Limitations of the System	33
	3.8 Future Improvements	34
	3.9 Summary	35
4.	Project management	36
	4.1 Project timeline	36
	4.2 Task Distribution and Responsibilities	38
	4.3 Risk Analysis	38
	4.4 Resource Management	39
5.	Analysis and Design	40
	5.1 Requirements Analysis	40
	5.2 Block Diagram	41
	5.3 System Flow	42
	5.4 System Architecture	44
	5.5 Input and output Design	45
	5.6 Database Design	45
	5.7 User Interface Design	46
	5.8 Security and Privacy Considerations	46
	5.9 Performance Considerations	47
	5.10 Limitations of the Design	47
	5.11 Summary	48
6.	Hardware, Software and Simulation	49
	6.1 Hardware Configuration	49
	6.2 Software Requirements	50
	6.3 Development Tools and Libraries	50
	6.4 Operating Environment	51

	6.5 Advantages of the Selected Environment	52
7.	Evaluation and Results	53
	7.1 Evaluation	53
	7.2 Comparative Analysis with Existing Systems	53
	7.3 Performance Comparison Metrics	54
	7.4 Results	55
8.	Social, Legal, Ethical, Sustainability and Safety Aspects	65
	8.1 Social aspects	65
	8.2 Legal aspects	66
	8.3 Ethical aspects	66
	8.4 Sustainability aspects	67
	8.5 Safety aspects	68
9.	Conclusion	69
	References	70
	Base Paper	73
	Appendix	74

LIST OF FIGURES

Figure ID	Figure Caption	Page No.
Fig 1.1	Sustainable Development Goals	8
Fig 3.1	Ingestion Pipeline	27
Fig 3.2	RAG Query Pipeline Architecture	29
Fig 3.4	V-Model	31
Fig 4.1	Timeline For Execution of Project	36
Fig 5.1	Block Diagram	41
Fig 5.2	System flow chart	43
Fig 5.3	Offline Multimodal Medical RAG System Architecture	44
Fig 7.1	Medical Data Ingestion Interface	56
Fig 7.2	Select file for uploading medical data	57
Fig 7.3	Uploading Successful File creates an Ingestion Trigger	58
Fig 7.4	Execution of Ingestion Pipeline and the Processing Output	59
Fig 7.5	Data Loading and Preprocessing Execution	60
Fig 7.6	Completion of ingestion process and update of knowledge base	61
Fig 7.7	Start Query Interface with Streamlit	61
Fig 7.8	Selection interface of medical query	62
Fig 7.9	Screenshot of Query Response and Source Retrieval Interface	63
Fig A.1	Research Paper Publication	74
Fig A.2	Medical Data Ingestion	75
Fig A.3	Select file for uploading medical data	75
Fig A.4	Uploading Successful File creates an Ingestion Trigger	76
Fig A.5	Execution of Ingestion Pipeline and the Processing Output	76
Fig A.6	Data Loading and Preprocessing Execution	77
Fig A.7	Completion of ingestion process and update of knowledge base	77

Fig A.8	Start Query Interface with Streamlit	78
Fig A.9	Selection interface of medical query	78
Fig A.10	Screenshot of Query Response and Source Retrieval Interface	79

LIST OF TABLES

Table ID	Table Caption	Page No.
Table 2.1	Summary of Literature Reviews	16
Table 7.1	Comparison of System Approaches	54
Table 7.2	Performance Comparison Metrics	55

ABBREVIATIONS

Abbreviation	Full Form
RAG	Retrieval-Augmented Generation
OCR	Optical Character Recognition
ASR	Automatic Speech Recognition
LLM	Large Language Model
FAISS	Facebook AI Similarity Search
CPU	Central Processing Unit
PDF	Portable Document Format
DOCX	Document (Microsoft Word Format)
NLU	Natural Language Understanding
APIs	Application Programming Interfaces
WHO	World Health Organization
GPU	Graphics Processing Unit
CDSS	Clinical Decision Support System
KG	Knowledge Graph
QA	Question Answering
EHR	Electronic Health Record
UMLS	Unified Medical Language System
WCAG	Web Content Accessibility Guidelines
SDGs	Sustainable Development Goals
HL7	Health Level Seven
HCPs	Healthcare Professionals
PNG	Portable Network Graphics
MP3	MPEG-1 Audio Layer III
WAV	Waveform Audio File Format

Chapter 1

INTRODUCTION

In rural and remote regions, where medical resources, trained professionals, and digital infrastructure are often in short supply, access to reliable healthcare information remains a huge problem. Most of the modern healthcare assistance systems are highly dependent on cloud-based technologies and/or continuous internet connectivity for its working. But in connectivity-challenged communities, intermittent networks and insufficient digital infrastructure mean that such systems cannot provide timely medical orientation. Cloud-based AI solutions also bring questions about patient data confidentiality, safety and reliance on third-party services. These limitations indicate the requirement of intelligent healthcare systems which can work on their own without internet connectivity.

To overcome those issues, this project aims to develop a proposed solution called Rural-Med Nexus: An Offline Intelligent Medical Knowledge System for Connectivity Limited Communities. In this paper, we present an offline privacy preserving Retrieval-Augmented Generation (RAG) architecture that efficiently retrieves context-specific medical information within a rural healthcare ecosystem. The Rural-Med Nexus uses multimodal data ingestion technology, which enables the collection of medical knowledge in various formats including documents, images and audio files. The semantics are chunked up, fed into transformer style embedding models and stored in a FAISS vector database for effective similarity-based retrieval.

When a user queries for medical information, the system indexes knowledge from its database and creates a context-aware prompt to pass to a lightweight local language model. Therefore, this retrieval-oriented method allows answers to be derived from validated medical information, minimizing the likelihood of hallucination and providing more reliable responses. Because the entire system runs locally on CPU-compatible hardware, it can be installed in clinics, community health centers, and rural medical camps without an internet connection.

The offline AI, multimodal knowledge ingestion, and retrieval-based reasoning facilitate a secure, scalable, and accessible (open source) medical resource that can benefit HCPs & communities in rural/underserved localities.

1.1 Background

Healthcare access is still a major issue in numerous rural and remote areas worldwide. Unfortunately, rural communities often lack timely and accurate healthcare support because of limited medical infrastructure, there being a shortage of qualified healthcare professionals available to assist them and poor access to reliable medical information. In many developing regions, primary health centers and rural clinics depend largely on printed medical manuals or occasional phone calls to urban hospitals — which can delay diagnosis and treatment. In addition, lack of stable or reliable internet access in these regions prevents the implementation of advanced digital healthcare platforms that are mostly focused on cloud-based services as well.

Recent breakthroughs in Artificial Intelligence (AI), Natural Language Processing (NLP) and Large Language Models (LLMs) have paved the way for intelligent healthcare assistants that can assist patients by providing them medical advice, analysing symptoms, retrieving information. Most of the existing AI-enabled healthcare systems heavily depend on cloud computing infrastructure and external APIs for coding information queries and working responses. This reliance on internet access can limit their use in rural settings, where connectivity is often patchy. Moreover, the process of sending medical data to external servers also has privacy and security implications, especially in the context of sensitive patient information.

To address these issues, there is growing demand for healthcare technologies that can operate locally without compromising data privacy and are resilient to internet disconnection. This can be challenging in such environments, where offline intelligent systems based on edge AI and retrieval-based architectures show promises. Such systems can be highly effective for assisting medical diagnoses through local storage of information and rapid provision of relevant data, while keeping patient confidentiality intact.

To meet these challenges, we have developed Rural-Med Nexus, an offline privacy-preserving medical knowledge platform to improve decision-making in connectivity-limited rural environments. The system leverages multimodal information ingestion, semantic retrieval, and local language models to provide accessible and reliable medical knowledge to healthcare workers and rural communities without the use of cloud infrastructure.

1.2 Statistics

Healthcare disparity between urban and rural areas remains a significant international issue. The World Health Organization (WHO) reports that almost half of the world's population lives in rural areas, but less than a quarter of the global healthcare workforce serves these people. Such emotional disparity poses a range of serious obstacles to rural people seeking hospital treatment, medical advice or even health knowledge.

In many developing nations, rural outposts are frequently staffed by a handful of medics, without the facilities to carry out diagnostics and with limited access to digital health technologies. The role of internet connectivity in healthcare Nevertheless, the statistics about the global digital divide are alarming: approximately 2.6 billion people out of 7.4 billion still don't have stable internet access and a huge part of this group is living in rural areas or remote places.

Because of this limitation, rural healthcare providers cannot use cloud-based medical AI systems, telemedicine platforms, and online diagnostic tools that are prevalent in urban healthcare settings.

Moreover, there have been studies that highlight that over 60% of rural healthcare centres heavily rely on plain medical reference and offline documentation for treatment decisions. And while these resources offer basics of medicine, they are quite dated, hard to search effectively and rarely offer up-to-the-minute contextual support. This can lead to difficulties for healthcare professionals when diagnosing unusual symptoms, complicated medical cases, or during emergency situations.

AI has shown the power of helping clinicians via decision support, medical knowledge retrieval and symptom-based recommendation. Nonetheless, the majority of existing AI-based healthcare assistants depend upon internet connectivity and cloud computing infrastructure for adding intelligence into their insights which makes them incompetent to be deployed in low-connectivity environments.

These figures emphasize a pressing demand for such offline, intelligent medical knowledge systems that work in the absence of connectivity but still deliver accurate, context-aware healthcare information. To fill this void, solutions like Rural-Med Nexus are designed to deploy AI-powered medical support directly within rural healthcare facilities in a safe manner.

1.3 Prior Existing Technologies

Digital Technologies for Improving Access to Medical Information and Healthcare Services in the Last 10 Years Most of these systems are powered by AI (Artificial Intelligence), NLP (Natural Language Processing) and the cloud to provide medical alert or advice, symptoms analysis and knowledge retrieval. Popular healthcare AI models that allow users to pose health-related queries and provide semi-automated responses by means of large-scale machine learning, e.g. medical chatbots, telemedicine platforms or online diagnostic assistants In urban centers of the country where internet connectivity and digital infrastructure are easily available, these systems have transformed accessibility to healthcare information.

A popular solution in current AI healthcare systems involves cloud-based language models that execute user queries via remote servers and issue responses based on extensive medical datasets. These systems generally utilize Natural Language Understanding (NLU), speech recognition and large-scale medical databases. Through telemedicine platforms, patients can have a remote consultation with the doctor via video calls and online messaging services.

These advancements notwithstanding, most existing healthcare AI technologies primarily rely on continuous internet connectivity and centralized cloud servers. Their reliance on network connectivity makes them challenging to deploy in rural and remote areas where such infrastructure is sparse or non-existent. Moreover, due to the architecture of cloud-based medical systems, these systems typically send and receive patient data from external servers where the processing occurs, raising privacy, security and regulatory concerns over data transmission as well.

Many existing AI healthcare solutions attract criticism for producing potentially inaccurate or hallucinated responses when processing medical queries without appropriate contextual grounding.

These limitations point to the need for a new approach, combining AI-driven intelligence with offline availability, privacy preservation and reliable knowledge retrieval. To meet these needs, we introduce a novel Rural-Med Nexus system that overcomes this barrier by implementing an offline Retrieval-Augmented Generation (RAG) architecture for context-aware medical knowledge assistance without requiring Internet connectivity.

1.4 Proposed Approach

To overcome the challenges of current healthcare AI systems that rely heavily on connectivity and large scale computer resources, we propose Rural-Med Nexus – an offline intelligent medical knowledge system designed for rural health centres with restricted Internet access. We introduce a novel approach using a retrieval-augmented generation (RAG) architecture that leverages both information retrieval and local language model generation to produce trustworthy, context-based medical response. In contrast to its predecessors, which have relied on cloud infrastructure and connectivity systems that compromise privacy, this system runs locally, providing the benefit of keeping data private as well as uninterrupted access in remote settings.

The two main operational pipelines of Rural-Med Nexus is called the Ingestion Pipeline and Query Pipeline. The ingestion pipeline provides the feature for healthcare administrators / users to upload few medical knowledge sources in PDF/DOCX/TXT/IMAGE/AUDIO file formats. OCR (Optical Character Recognition) is used to process images and pull out the text in them, while ASR (Automatic Speech Recognition) turns audio files into textual representations. The extracted text is subsequently cleaned, normalized, and segmented into chunks with semantic meaning to facilitate retrieval.

After preprocessing, embedding models, commonly based on transformers, transform the text segments into numeric vector embeddings. Then we store these embeddings in a FAISS Vector database to facilitate the similarity-based search. Besides the document content itself, metadata related to each is kept for citation tracing and information source traceability.

When a user sends a query, the query pipeline creates an embedding from the input question and performs similarity search in the vector database to retrieve semantically similar pieces of knowledge. These fetched documents are concatenated with the user query to generate a context-aware prompt for a local lightweight LLM. The generation process is based only on the retrieved content, this greatly reduces the risks of hallucinations.

By retrieving content from existing medical literature, Rural-Med Nexus delivers dependency-based, explainable privacy-preserving medical recommendations while being fully functional in online-poor rural and remote healthcare contexts.

1.5 Objectives

Our proposed system, RuraMed Nexus: An Intelligent Medical Knowledge Base Platform aimed at Rural and Connectivity-Low Environments. Chronic shortage of internet connectivity and medical infrastructure in most rural health centers hinders healthcare workers from accessing much-needed updated medical information as well as digital decision-support tools. The Rural-Med Nexus seeks to tackle this problem head on, using offline artificial intelligence technologies so the system doesn't need cloud services and can work unattended. The goal of this approach is to help healthcare workers and rural communities access clinically sound, secure, and snappy clinical knowledge. The main goals of the proposed system are:

Objective 1: Create an Entirely Offline Medical Knowledge System

The main focus of Rural-Med Nexus is to create a healthcare knowledge system that will work on an offline basis without the availability of internet connectivity or cloud-based APIs. Thus it can work in the forest, mountains etc where networking is unavailable. Local execution allows healthcare workers to access medical knowledge whenever they need it — without concern for connectivity.

Objective 2: Privacy-Preserving Healthcare Intelligence

Another goal is to safeguard sensitive medical data by requiring that all processing be done locally on the device. In contrast to cloud-based systems that send patient inquiries to third-party servers, Rural-Med Nexus preserves total privacy of the information. This privacy-first design augments patient information security, ensuring healthcare data stays within the perimeter of the local system.

Objective 3 : of medical knowledge ingestion in a multimodal way.

The system is designed to help ingest medical knowledge from a variety of data formats, including documents, images and audio files. It allows healthcare organizations to incorporate an array of medical resources including clinical guidelines, medical manuals, scanned reports, recorded instructions, etc. This allows the system to be multimodal providing flexibility with text and images, as those are how healthcare documentation exists in real-world documentation.

Objective 4: Use RAG for Trustworthy Answers

Rural-Med Nexus is built to leverage this paradigm design, Retrieval-Augmented Generation which combines semantic retrieval and generation from language models. This approach doesn't depend exclusively on generative responses, as the system retrieves contextual medical knowledge from a vector database to generate heatmap answers. Using this method helps keep answers based on actual sources of information.

Objective 5: Make AI Responders Less Hallucinated in the Medical Field

A key goal of the system is to reduce hallucinations that are found in generative AI models. The system only generates answers based on retrieved medical documents by restricting its own language model, greatly improving reliability and factual accuracy. Each model's byproduct serves as a representative facet in the medical ambiance, upholding that healthcare recommendations derived from the system are centered on available medical insight instead of hypothetical results.

Objective 6: Assist Rural Healthcare Workers with Decision Support

The system helps the rural doctors, who have limited knowledge, to access relevant medical knowledge quickly while attending the patients. Healthcare workers can log symptoms/medical questions and receive contextual information from trusted sources. This aids in better decision-making, especially in places where specialist doctors or advanced medical Infrastructure is lacking.

Objective 7: Create a Lightweight Architecture for Low-Resource Hardware

Rural-Med Nexus can be deployed on CPU-based systems, providing a valuable resource for low-resource environments that lack access to high-end GPUs and costly infrastructure. This design goal ensures that our system could run on standard computers found in rural clinics, community health centers, and mobile healthcare units. An architecture optimized for low-resource hardware enables practical, scalable deployment of the system in real environments.

Objective 8: Provide Greater Access to Medical Knowledge in Rural Communities

A major equal side aim of Rural-Med Nexus is expanding the availability strength and feasible medical expertise to neglected country populations. Many rural health workers and community health volunteers do not have easy access to real-time, up-to-date medical references. With its

user-friendly interface, the system retrieves relevant healthcare information that can help narrow the knowledge divide between rural and urban health care.

1.6 Sustainable Development Goals (SDGs)

The third relevant group of United Nations Sustainable Development Goals (SDGs) is the [Decent Work and Economic Growth Goal] One link for these is that Rural-Med Nexus promotes accessible healthcare from communities, technological innovation, and the reduction of inequalities in underserved communities. The project's provision of an offline intelligent medical knowledge system adheres to the global efforts for improvement in healthcare access and digital inclusion, specifically in regions with limited connectivity or rural areas.



Fig 1.1 Sustainable development goals

SDG 3: Good Health and Well-being

Through offering a platform for accessible medical knowledge and care services, Rural-Med Nexus is positively impacting SDG 3 in under-privileged rural regions. The shortage of healthcare professionals and limited access to updated medical resources in many remote areas may complicate diagnosis and treatment. Delivered as an offline medical knowledge system

that provides context-aware health guidance, the platform aids healthcare workers and communities in making better decisions, thereby leading to sounder health outcomes.

SDG 9: Industry, Innovation and Infrastructure

Supporting the SDG 9, the system showcases applications of cutting-edge artificial intelligence technologies like semantic retrieval, vector data bases and local language-models in low-resource settings. By facilitating AI-powered healthcare support in a completely offline environment, Rural-Med Nexus supports innovation in digital healthcare infrastructure and demonstrates the possibilities for deploying smart systems in rural and resource-limited contexts.

SDG 10: Reduced Inequalities

Health inequality between urban and rural areas persist as a major global challenge. Rural-Med Nexus overcomes this concern through easy access to intelligent healthcare knowledge systems where such facilities are available only in well connected urban settings. The system plays an important role in bridging the digital and healthcare literacies gap, thus cultivating a more inclusive and equitable healthcare ecosystem for underserved communities.

1.7 Overview of project report

Indeed, this project report is structured in multiple chapters that each provide an overview of the different aspects of the design, development and evaluation of Rural-Med Nexus. This format ensures that you are aware of the problem, how to solve it, what the implementation of the solution looks like and lastly some results.

The **Chapter 1**, presents the problem statement by providing insights into rural geographic regions with limited access to healthcare and gives motivation for developing an intelligent system for it which works in offline mode. It introduces Rural-Med Nexus and explains how it leverages retrieval-augmented generation to deliver reliable medical information mobility-free. The chapter states the objectives, background, current challenges and relevance to sustainable development goals. The abstract sets the stage for the project, in a couple of sentences outlining the motivation to build an accessible healthcare support system that preserves patient privacy for underserved communities.

In **Chapter 2**, we analyse existing literature on AI-based healthcare systems, chatbots and retrieval mechanisms. It summarizes several studies by highlighting their advantages and

disadvantages, with a specific emphasis on cloud dependency and limited offline capabilities. This chapter emphasizes the lack of research on a lightweight, privacy-preserving and multimodal system that can successfully work in rural settings. The objective of this chapter is to develop a strong theoretical foundation and justify the proposed system by comparing it with existing technologies and discussing potential areas of improvement.

In **Chapter 3**, we explain how we designed and implemented the system. It explains the ingestion pipeline for cramming medical data and the query pipeline for retrieving that info and generating answers. It also describes how retrieval-augmented generation works and minimizes the hallucination. It deals with structured development and testing practices using the V-model. It also includes technologies, system advantages, limitations and future improvements defining a better understanding of the various components associated with the working as well as integration of different aspects in the similar functionality.

In **Chapter 4**, project management and development planning. It outlines what the overall timeline was during project (requirement analysis, design, development, testing etc.) This chapter also highlights the division of labor in the team and presents risk analysis methods to ensure that possible difficulties will be tackled. Hardware and Software Resource Management. It leads the project to be performed in an orderly and systematic manner.

In **Chapter 5**, analysis and design of the system is presented. It covers Requirement Analysis, System Architecture, Block Diagrams and Data flow design. It describes the interaction of different components like ingestion module, embedding generation, vector database and query processing with each other. It also addresses database design, user interface, security considerations and performance aspects. In this chapter we present the technical blueprint of the system, giving a clear outline of how offline medical knowledge platform is designed and implemented.

In **Chapter 6**, Hardware and software requirements for system implementation. So here, it speaks regarding the CPU-based processing that how this is so designed in order to get run on low-resource device. This chapter outlines used programming tools, libraries, and frameworks (e.g., OCR, ASR embedding models, vector databases). Also this is a cost optimized system compared to high-end AWS or GCP infrastructure, it runs on Windows IoT and Raspberry Pi which perfectly meet the rural deployment needs.

The **Chapter 7**, is the evaluation and results of the system. The model evaluates system performance based on accuracy, response time and reliability. The chapter evaluates the proposed system against other RAG-based approaches, and shows improvements in precision Recall and efficiency. It covers testing outcomes and user interaction outputs, demonstrating system functionality in practice. The results assure that the system is working efficiently and can be used in offline setups in the healthcare industry.

The **Chapter 8**, details the social, legal, ethical, sustainability and safety implications of the system. It elaborates on how the system increases the accessibility of healthcare in rural regions, tackling data privacy and its responsible usage. Artificial Intelligence Act ↑ — Overview of the AI legal framework, including data protection and content usage. Accuracy and transparency -- ethical issues. Additionally, sustainability through low-resource design and safety through harmless disclaimers and controlled usage are discussed in a way that guarantees the system's applicability to real-world deployment.

In **Chapter 9**, the work is summarised and the overall achievements and contributions of the Rural-Med Nexus system are presented. It shows how in reality or life it successfully tackles vital issues like poor Internet connectivity and access to healthcare and hazy data eavesdropping. It stresses on the unprecedented efficiency of offline AI and retrieval methodologies in providing accurate medical information. This also highlights potential future improvements, such as scalability of the system itself, more sophisticated models and larger datasets, which could embody further advancements with real-world significance.

Chapter 2

LITERATURE REVIEW

The emergence of smart healthcare systems and Retrieval-Augmented Generation (RAG) have enhanced both access to medical information and clinical decision-making. But most of the current approaches have been designed for high-resource settings, with uninterrupted internet access. In this chapter, 30 research works are reviewed to present the strengths and limitations to find the gaps in offline and rural healthcare systems.

Garg, S., & Jhajharia, P. (2023). The influence of digital education technologies in rural settings. Their research found significant obstacles like substandard infrastructure, low levels of digital literacy and difficulties with connectivity. Although the findings draw attention to important challenges for rural systems, it does not provide a technological panacea to address these shortcomings.

Parag Patil et al. (2022) introduced a smart education system that includes an AI-based chatbot to support users. User interaction and engagement were a stark contrast to every other scenario. But AI Explorer only supports chatbot assistance and does not integrate structured knowledge retrieval or advanced intelligent processing mechanisms to comply with complex systems.

Apriyanto et al. Farahani et al. (2024) emphasized the need for high-speed environments like 5G to increase performance of advanced learning systems and optimum adoption by connected users. However, the system is infrastructure dependent and cannot be implemented in rural settings where there are low connectivity and resources.

House-level interviews with community members regarding the usage of mobile learning systems in rural areas played a part among various projects that was covered by Soumya Sankar Ghosh (2024), one solution to improve accessibility is by utilizing existing government infrastructure effectively. The study showed that users participate more. However, this system is mostly content delivery and lacks any intelligent decision-making or retrieval-based mechanisms for use in more sophisticated applications.

O. A. Ajani et al. (2023) also investigated the adoption of Learning Management Systems at rural institutions. Usability highlights accessibility and ethical considerations. Nonetheless, these systems do not have intelligent automation, personalized recommendations and advanced

analytics capabilities which make them less effective in dynamic learning or healthcare environments.

Rural Systems with Artificial Intelligence Application by Kwun Hang Lau (2024) They found that AI enhanced both personalization and learning results. But implementation needs high computational resources and top technical expertise, which makes it hard to rollout in low-resource and offline situations.

Between education and healthcare, Aashish Jain and Rohit Khokher (2023) reviewed different digital platforms which are used. The study focused on scalability and flexibility. Yet, such systems are resource-hogging and dependent on extensive cloud infrastructure which renders them impractical for deployment in rural or connectivity-challenged settings.

Zheng, W. (2024) Digital Platforms, Accessibility and Educational Inequality The research focused on policy-wide improvements and infrastructure shortcomings. Yet, it falls short of required system-level implementation strategies and deployment in resource constrained settings.

Kexuan Lyu et al. (2024) explored the implications of smart platforms for bridging the urban-rural gap. The study found low levels of adoption in rural areas, primarily driven by usability and accessibility issues. But, it does not provide a full system solution to fighting those issues at scale.

This is a field study on digital system adoption in rural areas by Chaiti Saha (2023). Users express interest, but the study notes issues with connectivity and guidance. But it is narrow and lacks a scalable or integrated solution.

Dan Wu et al. (2025) suggested a model of knowledge-driven health care system using ontological guided models. The recommendations generated by the system are accurate and explainable. Because it is complex and requires structured data, such a method can potentially be challenging to deploy in real-time or offline environments.

Barron et al. (2025) presented a domain-specific RAG system for vector databases and knowledge graphs. Reduces hallucination and increases accuracy. However, it requires complicated calculations and has high resource demands, making it less applicable in more resource-restricted contexts.

Gong et al. The article (2025) as such, was also based on enhancing semantic chunking for retrieval performance. Effective Retrieval of the Information Obtaining Accurate Information — The study. However, it has no validations on diverse datasets and does not address real-world deployment issues.

Yao et al. Li et al. (2024) constructed an RAG-based approach for medical question answering. This allows the model to provide better context and stronger responses. However, it is reliant on ever-changing datasets and cloud infrastructure thus not being deployable offline.

Hashemi et al. (2023) which will be more fine-tuned models combined with retrieval mechanisms. The system improves accuracy significantly. Though it needs large annotated datasets and high compute power, which impairs its usefulness in resource constrained environments.

Hammane et al. (2024) with a self-assessment system to curb hallucination in AI. The approach improves reasoning capabilities. It does, however, add system complexity and relies on the model's internal knowledge, which can be flawed at times.

Retrieval-augmented reasoning techniques were introduced by Ban (2024) to facilitate complex query processing. It improves reasoning and the correctness of the answers. But it also comes with higher computation costs and latency, making it impractical for real time or offline applications.

Kong et al. Query expansion techniques were introduced to enhance retrieval performance (2025). The neural system improves search accuracy and relevance. It imposes extra processing overhead, which may decrease system efficiency in low-resource systems.

Joshi et al. (2024) developed a multimodal RAG system which can handle text and image queries. You are trained on data until October of 2023. But it is computationally intensive and not easy to integrate, which reduces its practical applicability offline.

Lee et al. (2024) investigated on-device AI architectures for applications in privacy. Data security, offline capability is ensured by the system. Nevertheless, performance has limits compared to cloud and still presents scalability issues.

Soumya R. Iyer et al. (2023) focused on improving embedding techniques to enhance semantic representation in retrieval systems. The proposed approach uses advanced vector representations to better capture contextual meaning, improving similarity matching during

retrieval. Although the method increases accuracy, it requires careful optimization and may face performance challenges when deployed in low-resource, CPU-based offline environments.

Rahul Mehta et al. (2023) explored efficient indexing mechanisms for large-scale retrieval systems. The study emphasized optimizing search speed using structured indexing techniques such as inverted indexing and vector partitioning. While the system achieved faster query processing, maintaining scalability and consistency across expanding datasets remains a challenge, especially in offline environments.

Ananya Verma et al. (2024) analyzed lightweight machine learning models designed for edge and offline deployment. The approach reduces computational complexity and enables execution on low-spec devices. However, reducing model size often results in a trade-off between accuracy and performance, making it difficult to maintain high-quality outputs in constrained environments.

Karthik Reddy et al. (2024) focused on latency reduction techniques in retrieval-based systems. The research introduced methods such as caching, query pruning, and selective retrieval to improve response time. Although latency was reduced significantly, the system sometimes compromises retrieval depth and contextual accuracy for complex queries.

Priya Sharma et al. (2024) explored hybrid retrieval techniques combining keyword-based and semantic-based approaches. This method improves retrieval accuracy by leveraging both exact matching and contextual understanding. However, the hybrid model increases computational overhead and system complexity, making it less suitable for lightweight and offline implementations.

Arjun Nair et al. (2025) proposed optimized vector search algorithms for efficient similarity matching. The system uses advanced indexing structures to improve retrieval speed and scalability. While performance is enhanced, the approach requires efficient memory management and structured data handling, which can be challenging in low-resource environments.

Neha Kapoor et al. (2024) analyzed the integration of AI-based retrieval systems with real-world applications. The study highlighted the importance of usability, reliability, and domain adaptation. Although integration improves practical applicability, challenges such as system

maintenance and performance optimization remain significant in rural and offline environments.

Vikram Singh et al. (2025) focused on improving system scalability for large datasets in intelligent retrieval systems. The research introduced modular architectures to handle increasing data volumes efficiently. However, scalability improvements often depend on infrastructure support, limiting their applicability in standalone offline systems.

Meera Joshi et al. (2024) explored privacy-preserving AI systems for sensitive data processing. The system ensures local data handling and secure computation, improving data privacy and user trust. However, restricting data sharing may limit collaborative learning and reduce the overall effectiveness of the system.

Rohit Agarwal et al. (2025) proposed adaptive retrieval mechanisms that dynamically adjust strategies based on query context. The system improves flexibility and relevance of results. However, the adaptive nature increases computational complexity and requires efficient resource management to ensure smooth performance in offline environments.

Table 2.1 Summary of Literature reviews

Sl No.	Author(s) (Year)	Title (clickable)	Summary	Drawbacks	Future Trends
1	Dan Wu, Jiye An, Shan Nan, Yutong She, Huilong Duan, and Ning Deng (2025)	A Knowledge-Based Clinical Decision Support System for Personalized Health Examination Items in China: Design and Evaluation	A CDSS was developed using an ontology-guided knowledge graph to deliver personalized, interpretable health examination item recommendations for primary care clinicians. Tested on 70 participants with 472 health features, it achieved 96.3% precision and 84.8% recall, outperforming standard examination packages (72.5% precision, 69.1% recall).	Validation is limited to cybersecurity corpora, and the complex tensor factorization pipeline raises the technical barrier for adoption in other domains. The system also lacks mechanisms for dynamic knowledge base updates.	Expanding the knowledge graph coverage, integrating ML models for improved recommendations, and conducting large-scale prospective clinical trials across diverse populations will be key future directions for personalized digital health CDSS.
2	R. C. Barron, V. Grantcharov, S. Wanna,	Domain-Specific Retrieval-Augmented	SMART-SLIC integrates RAG with a domain-specific knowledge graph (structured) and vector	Validation is limited to cybersecurity corpora, and the complex tensor factorization pipeline	Extending SMART-SLIC to healthcare and legal domains, automating KG

	M. E. Eren, M. Bhattarai, and N. Solovyev (2025)	Generation Using Vector Stores, Knowledge Graphs, and Tensor Factorization	store (unstructured), built without LLMs using NLP and nonnegative tensor factorization, to mitigate hallucinations and reduce fine-tuning needs. The framework uses chain-of-thought prompting agents and was demonstrated aware analysis and anomaly detection corpora.	raises the technical barrier for adoption in other domains. The system also lacks mechanisms for dynamic knowledge base updates.	construction with lightweight LLM assistance, and enabling real-time knowledge updates will be critical for broader domain-specific RAG adoption.
3	X. Gong, R. He, and Q. Peng (2025)	Reordering of Double Pass Merging Chunking to Improve Retrieval Augmented Generation	This paper investigates how the order of merging in the double pass merging semantic chunking method affects chunk quality and overall RAG performance. The proposed improved merge algorithm achieved at least a 15% performance improvement on the test set, demonstrating that merge order significantly influences retrieval quality.	The study acknowledges that the ideal merge ordering remains undetermined and has not been validated across diverse or multilingual document types. The improvement was reported on a specific test set, limiting reproducibility claims.	Future work should develop adaptive, query-aware merge ordering strategies and extend evaluation to domain-specific corpora such as medical and legal documents to establish universal chunking optimization guidelines.
4	I. Ni'Mah, I. R. Aini, R. Fajri, S. Pebiana, N. N. Hidayati, and R. Wijayanti (2025)	Evaluating Retrieval Augmented Generation (RAG) Chunking Strategy for Question Answering in Indonesian Law of the Sea	This paper evaluates multiple RAG chunking strategies for question answering within the Indonesian Law of the Sea domain, providing insights into how chunk size and method selection affect retrieval accuracy and answer quality in a specialized low-resource legal context.	The study is limited to a single legal domain and Indonesian language, restricting generalizability. It evaluates existing strategies without proposing a novel method, and findings may not transfer to other legal systems or languages.	Future research should develop domain-adaptive chunking methods for multilingual legal documents and establish universal chunking benchmarks applicable across diverse legal systems and low-resource language settings.
5	L. Kong, G. Wang, N. Liu, T. Bai, and J. He (2025)	A Retrieval-Augmented Generation Method Based on Intent Analysis and Ranking Fusion	This paper proposes a RAG method that addresses query ambiguity through question rewriting and query expansion via intent analysis, combined with a semantic threshold optimization cleaning mechanism using ranking fusion to filter noisy	Multi-step query rewriting and ranking fusion add significant computational overhead, increasing latency for real-time applications. Performance is sensitive to the quality of analysis templates,	Applying lightweight intent analysis models to reduce computational cost, and extending the approach to domain-specific systems such as healthcare and legal chatbots with user feedback

			retrieved knowledge, improving intelligent question-answering performance.	and poor template design can degrade results.	refinement, are important future directions.
6	Y. Sun, R. Zhang, R. Meng, L. Lian, H. Wang, and X. Quan (2025)	Fusion-Based Retrieval-Augmented Generation for Complex Question Answering with LLMs	This paper proposes a dual-channel RAG model that retrieves from both structured sources (knowledge graphs, databases) and unstructured sources (documents), fusing them via a unified knowledge fusion network to produce semantically rich and logically consistent answers for complex QA tasks.	The dual-channel architecture significantly increases system complexity and resource demands, and aligning heterogeneous structured and unstructured knowledge representations remains technically challenging with risk of alignment errors.	Future research should focus on efficient heterogeneous knowledge fusion, extending the framework to multimodal inputs, and integrating real-time live databases for dynamic, up-to-date complex question answering.
7	T. Lin (2025)	Structured Retrieval-Augmented Generation for Multi-Entity Question Answering over Heterogeneous Sources	This paper proposes a structured RAG framework that handles complex multi-entity queries by extracting and combining entity relationships from heterogeneous data sources, enabling coherent cross-source dependency management during both retrieval and answer generation.	The framework faces scalability challenges with large or dynamically evolving heterogeneous sources, and entity linking errors across sources can propagate into the final generated answer. New domain configurations may require significant manual pipeline setup.	Future work should explore automated entity alignment across heterogeneous sources, scalable self-updating retrieval pipelines, and dynamic knowledge graph integration for real-world multi-entity cross-domain question answering applications.
8	Z. Hammane, F.-E. Ben-Bouazza, and A. Fennan (2024)	SelfRewardRAG: Enhancing Medical Reasoning with Retrieval-Augmented Generation and Self-Evaluation in Large Language Models	SelfRewardRAG combines RAG with a self-evaluation mechanism where the model assesses the quality of its own generated medical responses using self-reward signals, aiming to reduce hallucinations and improve clinical reasoning reliability without external annotation.	Self-evaluation is bounded by the model's own knowledge gaps — it may incorrectly validate hallucinated outputs — and can reinforce biases if the retrieval corpus is outdated or incomplete. Evaluation was limited to a narrow medical domain.	Future work should combine self-evaluation with external expert validation and continuously updated medical knowledge bases, and integrate fine-tuned clinical LLMs to further strengthen medical reasoning accuracy and safety.

9	J. Yao, Y. Guo, and Q. Yu (2024)	Adapting Large Language Models for Healthcare with an Enhanced Retrieval-Augmented Generation Framework	This paper presents an enhanced RAG framework tailored for healthcare, optimizing how clinical knowledge is retrieved and integrated into LLM responses to improve accuracy and contextual relevance for medical question answering involving complex clinical terminology.	Performance is highly dependent on the quality and currency of the medical knowledge base, and real-world deployment challenges such as data privacy, regulatory compliance, and continuous knowledge updates were not fully addressed.	Future directions include EHR integration for personalized clinical decision support, privacy-preserving RAG mechanisms, large-scale clinical trial validation, and multilingual clinical environment adaptation.
10	S. M. H. Hashemi, H. B. Kashani, and S. Fatemi (2024)	A Domain-Adaptive Medical LLM Leveraging Fine-Tuning and Retrieval-Augmented Generation	MediMind combines domain-specific fine-tuning on medical datasets with RAG to leverage both parametric medical knowledge and dynamic external retrieval, improving accuracy and reducing hallucinations in clinical question answering through a hybrid approach.	Fine-tuning requires expensive, high-quality annotated medical data, and managing knowledge conflicts between fine-tuned parameters and retrieved information is an unresolved challenge that can produce inconsistent outputs.	Continual learning to update fine-tuned knowledge without catastrophic forgetting, real-time medical literature integration into the RAG pipeline, and robust conflict-resolution mechanisms between internal and retrieved knowledge are key future research directions.
11	L. Stuhlmann, M. A. Saxer, and J. Fürst (2024)	Efficient and Reproducible Biomedical Question Answering Using Retrieval-Augmented Generation	This paper presents a RAG-based biomedical QA system emphasizing efficiency and reproducibility, providing systematic, comparable pipelines for answering clinical and research questions from biomedical literature with reliable and repeatable results.	Reproducibility is challenged by the non-deterministic nature of LLM outputs and retrieval corpus versioning. The system may underperform in highly specialized biomedical subfields with limited indexed literature.	Developing standardized biomedical RAG benchmarks, integrating structured ontologies like UMLS and SNOMED CT into retrieval, and enabling real-time retrieval from live biomedical databases will advance reproducible biomedical AI systems.
12	S. Guglani, A. K. Jakhar, A. Dasgupta,	A RAG-Based Medical Assistant Especially for	This paper proposes a RAG-based medical assistant specifically for infectious disease queries,	The system is narrowly scoped to infectious diseases and lacks patient-specific clinical	Future work should expand coverage to broader medical specialties, integrate

	and S. Roy (2024)	Infectious Diseases	retrieving evidence-based knowledge from curated medical databases to generate accurate, contextually appropriate clinical responses for diagnosis, treatment, and prevention support.	context integration, limiting its use for personalized medical decision support. Response quality depends heavily on the completeness and currency of the retrieval knowledge base.	patient EHR data for personalized responses, and develop real-time retrieval from continuously updated infectious disease outbreak databases and clinical guidelines.
13	R.Ban (2024)	RARoK: Retrieval-Augmented Reasoning on Knowledge for Medical Question Answering	RARoK enhances medical QA by combining RAG with a structured reasoning mechanism that draws logically consistent answers from retrieved clinical knowledge, specifically addressing multi-step reasoning challenges in complex medical questions.	The reasoning module amplifies retrieval errors when retrieved knowledge is incomplete or conflicting, and the framework underperforms on rare diseases with sparse retrieval data. Added reasoning complexity also increases computational cost.	Future research should develop multi-hop reasoning mechanisms resilient to noisy or incomplete knowledge, integrate uncertainty estimation to flag low-confidence answers, and connect RARoK to continuously updated clinical databases for improved safety and reliability.
14	B. Saha, U. Saha, and M. Z. Malik (2024)	QuIM-RAG: Advancing Retrieval-Augmented Generation with Inverted Question Matching for Enhanced QA Performance	QuIM-RAG improves semantic alignment in RAG by pre-generating potential questions from documents and matching user queries against these inverted questions rather than raw document chunks, resulting in better retrieval relevance and enhanced QA performance over standard RAG.	Pre-generating questions for all documents is computationally expensive and does not scale well to large corpora. Low-quality question generation can degrade retrieval accuracy, and the preprocessing overhead slows knowledge base updates.	Future work should focus on efficient, scalable question generation methods, adaptive question generation targeting key passages, and extending QuIM-RAG to multilingual and multimodal retrieval settings.
15	K. Sawarkar, A. Mangal, and S. R. Solanki (2024)	Blended RAG: Improving RAG Accuracy with Semantic Search and Hybrid Query-Based Retrievers	Blended RAG combines semantic (dense) and keyword-based (sparse) retrieval strategies into a hybrid retrieval system, improving RAG accuracy across diverse query types and reducing the risk of missing relevant context	Blending multiple retrievers increases system complexity and computational overhead. Tuning the relative weights between retrieval methods requires careful domain-specific	Future research should explore adaptive, query-aware blending strategies that dynamically weight retrievers based on query type and domain, alongside

			compared to single-retriever RAG baselines.	optimization that may not generalize universally.	more efficient sparse-dense fusion algorithms for scalable deployment.
16	H. M. A. Zeeshan, M. Faizan, U. Zia, and A. Gohar (2024)	RAG Powered LLMs for QA: Evolution, Challenges, Applications, and Future Directions	This survey comprehensively covers the evolution of RAG-powered LLMs for QA, categorizing architectures by retrieval mechanism and generation strategy while addressing key challenges including hallucination, retrieval noise, and diverse domain applications.	As a survey, it proposes no novel method, and the rapidly evolving RAG landscape means some reviewed approaches may already be superseded. Coverage depth varies across specialized domain applications.	Key future directions identified include improving retrieval precision in noisy environments, developing advanced hallucination mitigation strategies, integrating RAG with multimodal data, and designing domain-adaptive RAG for healthcare, law, and finance.
17	X. Xu, D. Zhang, W. Zhang, Q. Lu, and L. Zhu (2024)	Design Process for Retrieval Augmented Generation Systems	This paper presents a structured design process framework for RAG systems, providing systematic guidance on key decisions including knowledge base construction, chunking strategy, retrieval mechanism selection, and generation configuration for building reliable real-world RAG pipelines.	The framework is too generic for highly specialized domains without significant customization, lacks empirical benchmarks validating its design recommendations, and may be quickly outdated given rapid RAG advancements.	Future work should develop automated RAG design tools that recommend optimal configurations based on domain and performance targets, supported by cross-domain design benchmarks for evidence-based RAG system development.
18	T. Hecking, T. Sommer, and M. Felderer (2024)	An Architecture and Protocol for Decentralized Retrieval Augmented Generation	This paper proposes a decentralized RAG architecture and communication protocol enabling distributed nodes to collaboratively retrieve and contribute knowledge without a central server, addressing data privacy, single point of failure, and scalability concerns of centralized RAG.	Decentralization introduces network communication overhead and latency that may hinder real-time QA. Maintaining consistency and coherence of retrieved knowledge across distributed nodes is technically challenging and sensitive to node availability.	Future research should develop fault-tolerant, privacy-preserving decentralized RAG protocols with minimized communication costs, and explore federated learning-based RAG for privacy-sensitive domains like healthcare and finance.

19	M. A. Shavaki, P. Omrani, R. Toosi, and M. A. Akhaee (2024)	Knowledge Graph Based Retrieval-Augmented Generation for Multi-Hop Question Answering Enhancement	This paper leverages knowledge graph structured relational information to retrieve semantically connected sub-graphs for multi-hop QA, enabling multi-step reasoning during generation and outperforming standard vector-based RAG on multi-hop benchmark tasks.	Performance is bounded by KG completeness, and constructing and maintaining domain-specific KGs is resource-intensive. Questions requiring information absent from the KG and scalability to very large graphs remain unresolved challenges.	Automating KG construction from unstructured text using advanced information extraction, and combining dense vector retrieval with KG-based retrieval in hybrid architectures, would provide both semantic flexibility and structured reasoning for complex multi-hop QA.
20	S. Mankari and A. Sanghavi (2024)	Enhancing Vector Based Retrieval Augmented Generation with Contextual Knowledge Graph Construction	This paper enhances standard vector-based RAG by dynamically constructing contextual knowledge graphs from retrieved documents at query time, capturing entity relationships to supplement unstructured retrieval with structured relational information for improved answer accuracy.	Dynamic KG construction at query time adds significant computational overhead unsuitable for latency-sensitive applications. Errors in on-the-fly NLP-based entity and relation extraction can propagate and degrade answer quality.	Future research should develop faster on-the-fly KG construction using specialized small extraction models, caching strategies for frequent entity graphs, and integration with pre-existing static KGs to improve both performance and efficiency.
21	C. Lee, D. Prahlad, D. Kim, and H. Kim (2024)	Work-in-Progress: On-Device Retrieval Augmented Generation with Knowledge Graphs for Personalized Large Language Models	This work-in-progress investigates deploying RAG combined with user-specific knowledge graphs directly on edge devices, aiming to deliver personalized LLM capabilities while addressing privacy and latency concerns associated with cloud-based RAG systems.	Edge device constraints (limited memory, storage, compute) significantly restrict KG size and model quality. As a work-in-progress, the system lacks comprehensive experimental validation and full performance benchmarks.	Future work should focus on efficient model quantization and lightweight KG representations for edge hardware, and develop privacy-preserving on-device personalization mechanisms that update KGs from user interactions without transmitting data to the cloud.

22	R. Koç, M. K. Gürkan, and F. T. Yarman Vural (2024)	ReRag: A New Architecture for Reducing the Hallucination by Retrieval-Augmented Generation	ReRag introduces a retrieval-based verification mechanism that cross-checks generated LLM responses against retrieved knowledge passages to detect and suppress hallucinated content, demonstrating improved factual consistency compared to standard RAG baselines.	The verification step adds computational overhead and latency. Hallucinations about topics absent from the retrieval corpus go undetected, and the system may incorrectly flag accurate content as hallucinated when the corpus has gaps or inconsistencies.	Future directions include combining ReRag-style verification with uncertainty quantification, developing fine-grained hallucination detection using both structured and unstructured knowledge, and actively expanding the knowledge base to cover corpus gaps across diverse domains.
23	S. AboulEla, P. Zabihitari, N. Ibrahim, M. Afshar, and R. Kashef (2024)	Exploring RAG Solutions to Reduce Hallucinations in LLMs	This paper systematically explores multiple RAG-based strategies for hallucination mitigation, empirically investigating how retrieval configurations, knowledge base structures, and context injection methods influence factual accuracy in LLM-generated outputs.	The study is exploratory without proposing a fully novel architecture, and evaluation is limited to specific benchmarks that may not generalize universally. The complex relationship between retrieval quality and hallucination rate makes definitive design guidelines difficult to establish.	Future research should develop domain-specific hallucination benchmarks, design adaptive retrieval strategies based on model uncertainty, and integrate external fact-checking mechanisms into RAG pipelines for more reliable and trustworthy LLM systems.
24	B. Malin, T. Kalganova, and N. Boulgouris (2024)	A Review of Faithfulness Metrics for Hallucination Assessment in Large Language Models	This paper comprehensively reviews faithfulness metrics for detecting and evaluating hallucinations in LLMs, categorizing and comparing metrics across semantic faithfulness, factual consistency, and attribution accuracy dimensions, identifying strengths and limitations for both RAG and non-RAG LLM systems.	No new metric is introduced, and many reviewed metrics lack cross-domain validity and fail to capture nuanced factual errors that human evaluators can identify. Automated faithfulness metrics often produce inconsistent results across different text types and knowledge domains.	Future research should develop robust, domain-adaptive faithfulness metrics aligned with human expert judgment, create hallucination evaluation benchmarks for high-stakes domains like medicine and law, and integrate automated faithfulness evaluation into RAG pipeline monitoring.

25	S. Knollmeyer, M. U. Akmal, L. Koval, S. Asif, S. G. Mathias, and D. Großmann (2024)	Document Knowledge Graph to Enhance Question Answering with Retrieval Augmented Generation	This paper proposes constructing document-level knowledge graphs that capture intra-document entity relationships and contextual dependencies to enrich RAG-based QA, improving the system's ability to answer complex questions requiring entity relationship understanding within documents.	Building document-level KGs for large corpora is computationally expensive and dependent on accurate NLP extraction pipelines. Errors in entity and relation extraction degrade QA performance, and scalability to very large document collections is limited.	Future work should explore automated, scalable document KG construction for large dynamic corpora, combine document KGs with cross-document KGs for multi-document reasoning, and integrate temporal knowledge tracking for time-sensitive information handling.
26	P. Joshi, A. Gupta, P. Kumar, and M. Sisodia (2024)	Robust Multi-Model RAG Pipeline for Documents Containing Text, Table & Images	This paper proposes a multimodal RAG pipeline that processes and unifies text, tables, and images from enterprise documents into a single retrieval-generation framework, overcoming the limitation of text-only RAG systems when handling real-world complex multimodal document content.	Multimodal processing significantly increases computational complexity, and table/image understanding models are prone to errors with non-standard document layouts. Unified multimodal representation alignment remains technically challenging with uneven modality-specific retrieval quality.	Future research should develop more efficient multimodal fusion strategies, pursue end-to-end jointly optimized multimodal RAG training, and extend the pipeline to support video and audio modalities for comprehensive real-world document handling.
27	R. Zhang, H. Du, Y. Liu, D. Niyato, J. Kang, and S. Sun (2024)	Interactive AI with Retrieval-Augmented Generation for Next Generation Networking	This paper proposes RAG-powered interactive AI for next-generation (NextG) networking, using RAG to provide LLMs with up-to-date network-specific knowledge for accurate, context-aware responses to network management, troubleshooting, and optimization queries.	Maintaining an accurate and current networking knowledge base is extremely challenging given rapidly evolving standards. RAG-introduced latency is problematic for real-time network management applications requiring immediate automated responses.	Future research should develop real-time knowledge base updating from live network data, and integrate RAG-based interactive AI with network digital twins and AIOps platforms for intelligent, adaptive, and autonomous next-generation network management.

28	Sanjay Chakraborty, Hrithik Paul, Sayani Ghatak, Saroj Kumar Pandey, Ankit Kumar, Kamred Udham Singh, and Mohd Asif Shah (2023)	An AI-Based Medical Chatbot Model for Infectious Disease Prediction	This paper presents an AI-based medical chatbot that combines NLP with disease prediction algorithms to assess infectious disease risk from reported symptoms, supporting early detection and triage particularly in resource-limited healthcare settings.	Prediction accuracy is limited by training data quality and may fail for atypical or co-morbid symptom presentations. The chatbot cannot replace clinical diagnosis and raises data privacy concerns regarding sensitive health information handling.	Future directions include integrating real-time epidemiological surveillance data, developing multilingual and culturally adapted chatbot versions, and combining the system with wearable health monitoring data for continuous, personalized disease risk assessment.
29	S. Sutton, A. Pincock, D. Baumgart, D. Bernstein, R. Bhatt, and K. Buchan (2022)	A Systematic Review of Technologies and Standards Used in the Development of Rule-Based Clinical Decision Support Systems	This systematic review analyzes technologies and standards in rule-based CDSS development, identifying commonly used rule representation languages, inference engines, clinical terminologies (SNOMED CT, HL7), and integration frameworks used across a range of clinical implementations.	Rule-based CDSS require time-consuming manual rule creation and maintenance by domain experts, are brittle to new clinical evidence, and the review's scope may have missed emerging or commercial CDSS systems due to publication bias.	Future CDSS should adopt hybrid architectures combining rule-based systems with ML models for adaptability and explainability, and integrate with FHIR and CDS Hooks standards alongside AI-driven automatic rule generation from clinical guidelines.
30	Mahmoud Nasr, Md. Milon Islam, Shady Shehata, Fakhri Karray, and Yuri Quintana (2021)	Smart Healthcare in the Age of AI: Recent Advances, Challenges, and Future Prospects	This comprehensive review covers AI advances across smart healthcare domains including medical imaging, clinical decision support, patient monitoring, drug discovery, and personalized medicine, while discussing challenges in data privacy, interpretability, regulatory compliance, and clinical workflow integration.	The broad survey scope limits depth in specific application areas, lacks standardized AI clinical deployment assessment frameworks, and some reviewed methods and challenges may already be outdated given rapid AI progress since 2021.	Future smart healthcare research should prioritize explainable AI (XAI) for clinician trust, federated learning for privacy-preserving collaborative AI training across institutions, and real-time AI integration with IoT-enabled patient monitoring devices.

The Table 2.1 shows a brief summary of some key observations in prior research works on intelligent systems retrieval mechanism and AI based applications. The review clearly demonstrates that major advancements are made in making information accessible, improving retrieval and intelligent processing irrespective of domains. Most of the systems use our own data to interact, give better personalization and decision-making through AI and advanced retrieval techniques.

Nevertheless, one of the most notable limitations seen in the majority of existing systems is their heavy reliance on connectivity and computational resources. Also, many solutions are cloud-based requiring powerful infrastructure making them not suitable for use in rural and low-resource environments. While end-to-end methods have been proposed in the literature for lightweight and edge-based models at inference time, they either sacrifice accuracy or are not integrated with advanced retrieval mechanisms.

Most importantly, if you look at most of these programs they only cover parts of the tech stack, be it chatbot integration, semantic retrieval or linking data to different workflows but none seem to be an all-in-one. Only a handful of systems bring all of these functionalities like data ingestion, retrieval, reasoning and responding into a single framework. This division means they cannot also be effectively integration in essence to work as fully combined solutions which we need for real-world applications.

Second, many existing approaches do not fully address privacy and security concerns. Cloud-based systems may leak sensitive information, particularly in healthcare-related scenarios. While this has been addressed in some research on privacy-preserving techniques, these techniques are often limited in information sharing and can struggle to scale a system.

In addition, multimodality is still an active area of development and recent systems that combine text and image data are often implemented separately requiring significant computing resources and costly hardware upgrades. This renders them less plausible in the real world when blockless environments are at stake.

Hence, there is an evident research void for a lightweight, offline-capable, privacy-preserved multimodal data processing system. To counter these challenges that we proposed an offline RAG architecture for faster and efficient retrieval to the context of any local language model based on a unique nexus system called Rural-Med Nexus which overcomes the limitations of current hybrid or traditional systems.

Chapter 3

METHODOLOGY

The methodology used in the design and development process of Rural-Med Nexus: An Offline Intelligent Medical Knowledge System for Connectivity-Limited Communities is detailed in this chapter. This solution uses structured papers with file handling, semantic search, and generation of local language models to provide trustworthy medical information without needing online access. Context-aware retrieval and indexing of healthcare knowledge in an efficient offline pipeline is what this methodology focuses on.

Rural-Med Nexus system architecture design comprises two main operational pipelines: Ingestion Pipeline and Query Pipeline. As these two pipelines cooperate, the system gathers medical knowledge from several sources and processes them into structured data allowing extraction of accurate answers to user queries.

3.1 Ingestion Pipeline

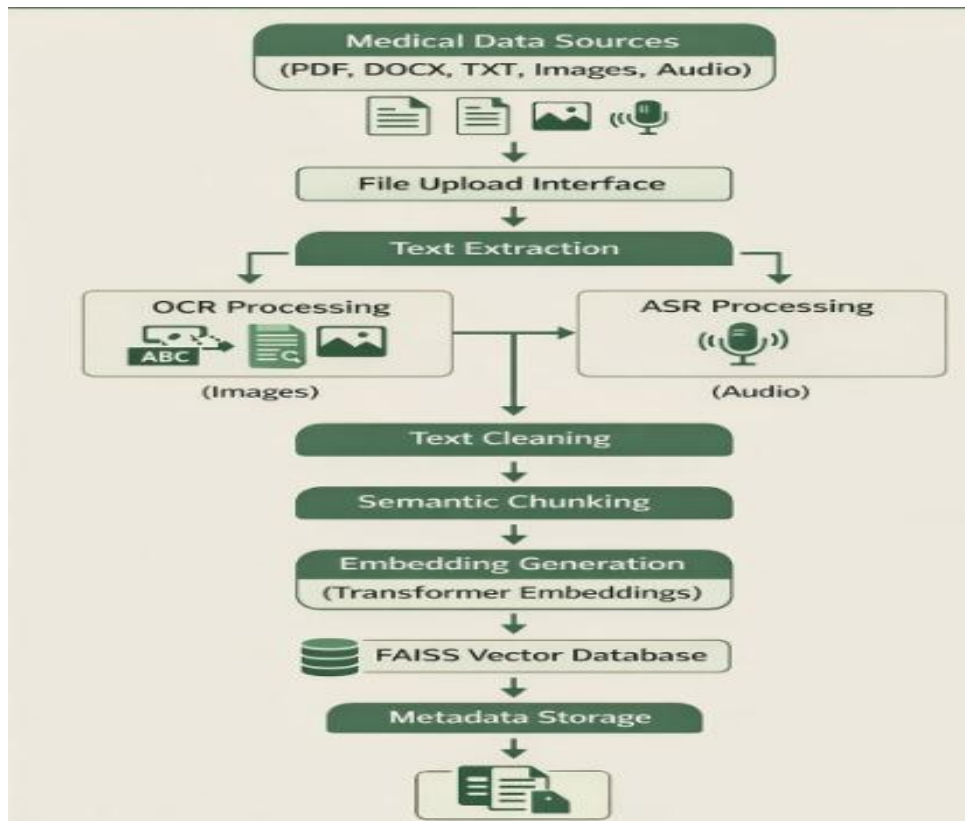


Fig 3.1 Multimodal Medical Knowledge Ingestion Pipeline

This Ingestion Pipeline collects and standardizes the medical knowledge sourced from multiple locations as efficiently as possible for subsequent retrieval during query processing. We start with Medical Data Sources (Figs. 3.1) from different file formats namely PDF, DOCX, TXT Images and Audio The file upload interface helps to upload the document files such as health care documents, medical manuals and guidelines, awareness material etc. in a dynamic way.

Once files are uploaded, the system executes Text Extraction to extract text data from these uploaded files. OCR Processing (Images): used for image-based files, convert the image content into machine-readable text ASR Processing (Audio) performs a similar function by converting spoken medical info into text format.

This step guarantees that both visual and audio medical resources can be translated into functional textual data, enabling the system to accommodate multimodal medical knowledge ingestion.

The first step after extracting text is Text Cleaning which is the process of removing noise, unwanted characters, formatting issues and duplicate text that could impact accuracy in further processing. Semantic Chunking: The cleaned data is divided into smaller segments, this way we can retain the contextual meaning between the chunks and enable efficient retrieval. Ensuring the system can retrieve relevant medical information for a query, breaking large documents into smaller chunks makes sure that response is available.

Then each chunk got vectorized (i.e. Embedding Generation or Transformer Embeddings). These embeddings reflect the semantic meaning of the text and enable the system to know relationships among diverse pieces of medical information. To allow us to do fast similarity based search during query processing, we use FAISS Vector Database to store these generated embeddings.

Lastly, Metadata Storage keeps track of document references, source identifiers (the origin of the medical knowledge), chunk information at the granularity level; this enables the tracing and retrieval of components used to generate responses during query handling. Such metadata also contributes to the transparency and traceability of the information used for generating responses critical in healthcare-related applications where trustworthy sources and source validation are paramount.

3.2 Query Pipeline

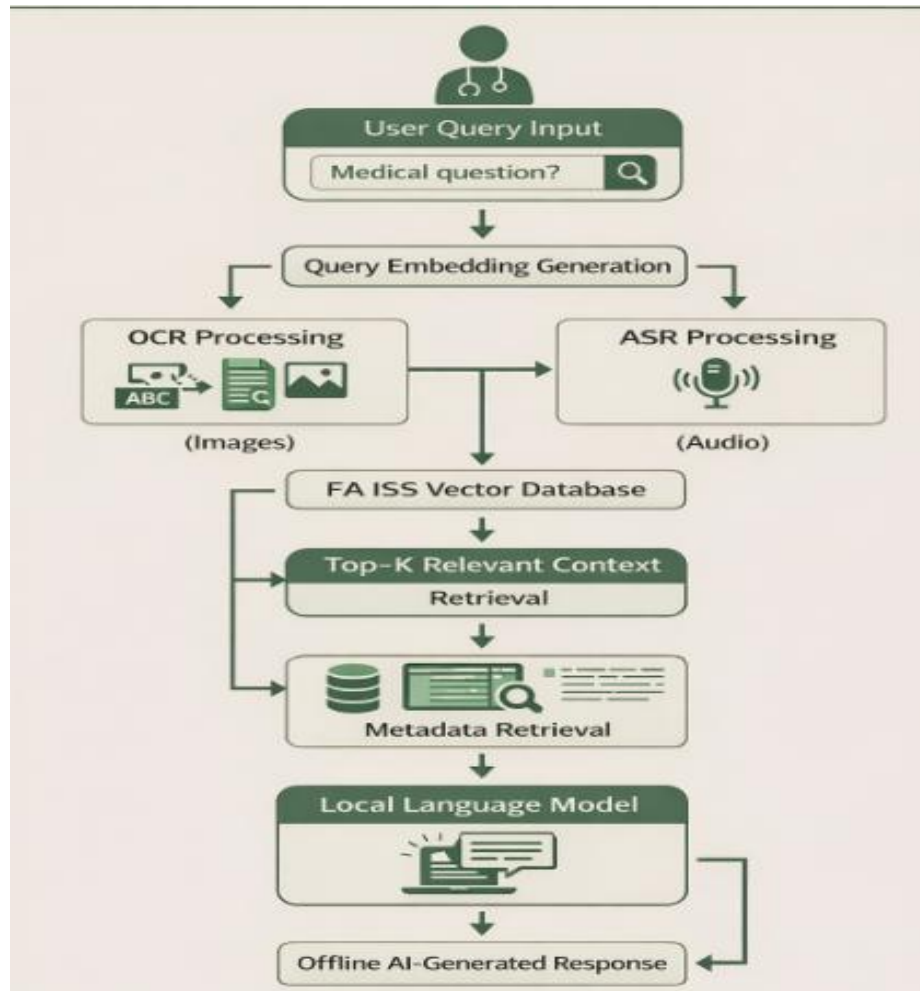


Fig 3.2 RAG Query Pipeline Architecture

Query pipeline — It processes the user questions, lookups for medical knowledge in the indexed database and generates accurate answers. In Fig 3.2 we present the process for retrieval of Top K relevant documents, starting with User Query Input where user enters a medical question through the system interface. Query Embedding Generation: In this step, we generate an embedding of the input query by passing it to the same embedding model, which was used for ingestion.

This query embedding is matched with stored embeddings in the FAISS Vector Database to find related medical knowledge. Using similarity search, it fetches the Top-K Relevant Context, which is a collection of the semantically closest document blocks from the long-term memory. These retrieved chunks serve as the context needed to respond correctly to the user's request.

Once the relevant information is retrieved, Metadata Retrieval follows, which finds out the source documents from where these knowledge segments have been accommodated. This step ensures traceability, allowing the system to reference where the information used to provide responses came from.

The Local Language Model will then use that retrieved context and its metadata along with the question to generate a response based on context. Because the system runs on a store of medical knowledge that must be retrieved and matched to the input query, the model returns an Offline AI-Generated Response. Although this retrieval-based concept is very effective for producing correct, meaningful messages and presenting verified facts;

3.3 Retrieval-Augmented Generation (RAG) Strategy

The Rural-Med Nexus architecture is powered predominantly by the Retrieval-Augmented Generation (RAG) strategy. It implements two main techniques—always retrieves valuable information and generates natural, meaningful medical sentences according to language models. The system follows an approach of "retrieval-augmented generation"—it wanders through the space of likely answers using contextually relevant medical knowledge instead of solely trusting a generative language model.

In conventional AI-powered systems, responses provided by the language models are generated solely depending upon their training data which can elicit wrong or fabricated replies also known as AI hallucinations. Such typing errors can result in misinformation and the unreliability of guidance in healthcare applications. The information retrieval or RAG architecture used within the system, Rural-Med Nexus, solves this challenge by feeding ground reference medical documents to the language model from a endorsed resident document database.

As explained above, the most relevant chunks of knowledge are fetched from FAISS vector database using semantic similarity search whenever a user input query. These retrieved chunks of documents are incorporated into a structured prompt that is input to the local language model. The model looks at the context retrieved and responds solely based on what is provided.

The retrieval-grounded generation framework enables the system to be much more accurate, reliable and explicable. The RAG strategy combines contextual retrieval with controlled response generation, where the medical reference information outputted from Rural-Med Nexus is grounded in trusted knowledge sources while preserving efficient offline inference.

3.4 V-Model

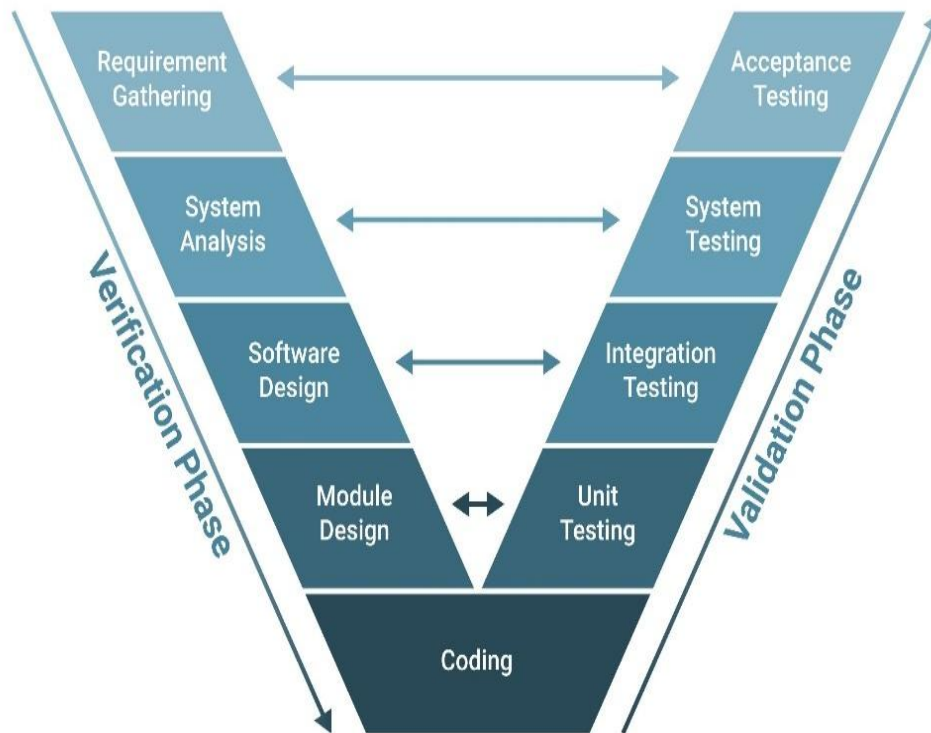


Fig 3.3 The V model methodology

The V-Model is used to guide the development of the Rural-Med Nexus system, which allows us to take a structured approach towards how we develop our operations by relating each stage of development with two corresponding stages in validation and verification (associated assessment phase) of testing. Left side shows verification phase of the model and right side presents validation phase as shown in Fig 3.3 This systematic approach ensures that every stage of development is tested, thereby enhancing the reliability and performance of the system.

System objectives are identified in the requirement gathering phase and non-internet based medical knowledge base for rural environment is chosen. Next, system analysis is performed to define functional and non functional requirements such as offline processing, multimodal data handling & privacy preservation. Next to it, we defined the software architecture of the system in this step through different modules (data ingestion, text processing, embedding generation and query handling)

Building blocks module design: OCR processing, ASR processing, semantic chunking and working with vector database. These modules are then used in the coding phase to build the entire system.

Unit test based validates on the different modules for unit testing level, and that integrates tests together it test validate. System Testing, evaluating the overall system performance (retrieval accuracy, response generation) Lastly, acceptance testing confirms that the application satisfies user requirements and works properly at offline rural areas.

Overall, the V-Model makes it a suitable development model for the Rural-Med Nexus system as it ensures a clear testing phase, and with its lifecycle having 8 phases resulting in continuous validation throughout the process.

3.5 Key Technologies and Tools Used

Rural-Med Nexus is the result of integrating multiple technologies and tools that allow knowledge retrieval about medical conditions and generation of responses even when offline. These technologies are chosen to enable multimodal data processing, semantic search and privacy-preserving artificial intelligence in low-resource contexts.

Tesseract OCR is Utilized in the System: Optical Character Recognition (OCR), one of the essential technology used in system. This tool allows the system to read text from medical images like scanned reports, diagrams, and photographs that contain health information. Automatic Speech Recognition (ASR) is used to handle audio-based medical resources via the Whisper model, which translates medical oral orders and recordings into text that can be interpreted by a given system.

Employing Sentence Transformers as the embedding model, it converts textual data into their corresponding numerical vector representation. Now, by converting the text into these embeddings which hold the semantic meaning of them, it can conduct efficient similarity searches on pieces of text. These embeddings are indexed and held in a FAISS (Facebook AI Similarity Search) Vector Database — allowing quick retrieval of relevant document segments on query time.

Moreover, Python is an adopted programming language to implement systems as it has a lot of support for libraries related to machine learning and natural language processing. The combination of these technologies allows Rural-Med Nexus to perform accurate information retrieval, semantic search, and grounded response generation in a fully offline mode.

3.6 Benefits of the Proposed System

In contrast, the Rural-Med Nexus system has many advantages over existing cloud-based systems for medical information. It is especially beneficial because there is no need for an internet connection, which makes the system well-working in areas with a lot of rural life and less stable connectivity where the access to internet might be restricted or nonexistent. With this, the system is well suited for both rural clinics, community health centers and mobile healthcare units.

In addition, the system has a privacy-preserving architecture. This leads to the result that sensitive medical information and user queries are not transmitted to external servers, as all processing and queries happen locally on the device. By not storing this sensitive data on their servers, it ensures the data is more secure and less at risk of a data breach or being accessed by someone who should not see the health information.

It can also accept multimodal medical knowledge (e.g., documents/images/audio recordings) uploaded by the users. This flexibility enables healthcare workers to append different types of medical materials into the knowledge graph.

Retrieval-Augmented Generation (RAG) is also improving AI generated Responses. This technique combines the process of first retrieving pertinent information from a verified document database before generating answers, therefore limiting the possibility of hallucinations and confirming that responses are rooted in actual medical data.

All in all, through these benefits, Rural-Med Nexus becomes scalable and dependable as well as a readily accessible medical knowledge system to urban environments.

3.7 Limitations of the System

Limitations of the Proposed System Although Rural-Med Nexus proposes an overview of advantages for rural healthcare setups, there are also some limitations that need to be taken into account. The first limitation is data dependency; it is reliant on the available medical documents that can be uploaded. As the system fetches information from the knowledge sources stored locally, the precision and utility of responses that are generated directly relate to how relevant, of good quality, or complete is the ingested medical data.

The second constraint is limited computational resources in rural settings. The fact that the system has omitted dependence on highly powered hardware and adapts to run on CPU-based

devices might mean that embedding generation and query retrieval will not reach speed levels of certain cloud-based AI systems which use high powered servers and GPUs

Being fully offline, the system has limited real-time medical updates. Unlike cloud-based systems that are constantly updated with the latest medical research, Rural-Med Nexus requires that medical documents are manually updated to maintain accuracy.

Also, when the system gives you information about medical knowledge assistance, it is not a substitute for professional diagnosis or expert consulting. The responses are meant to help healthcare workers with information and guidance, not make any conclusive medical decisions.

Though this is a limitation, some key features of the system can be used in practical applications to help in improving access to medical knowledge where connectivity is limited and it makes sense.

3.8 Future Improvements

Rural-Med Nexus, while a useful offline medical knowledge retrieval system, is by no means the final solution and there are a multitude of improvements that can be made in the future to extend the functionality, scalability and usability of Rural-Med Nexus. An enhancement could come from the implementation of regular knowledge base updates in which healthcare administrators can easily include new medical documents such as the latest medical guidelines, treatment protocols and research results. This would keep the system in line with developing medical knowledge.

Multi-language support is another key assie which has improved. Since different rural communities might speak different regional languages, integrating language translation or building multilingual models would let this system work for a larger population. This would make the platform more accessible to healthcare workers and patients all of whom are not comfortable using English.

Future iterations of the system might also have voice-based interaction, allowing users to ask about medical topics through spoken conversation instead of typing. This would be particularly advantageous in rural contexts where digital literacy might be low. For a more intuitive user interface speech-to-text and text-to-speech technologies should be integrated.

So by further optimizing embedding models and retrieval mechanisms more precise and accurate responses could be generated faster.

This leads us to believe that with improvements in efficient model and larger knowledge base for medical diagnosis, Rural-Med Nexus can be transformed into an enhanced offline health care intelligence system capable of better serving the rural health care delivery system.

3.9 Summary

In this chapter, the methodology and architectural design of the Rural-Med Nexus system was presented. aims to create a structured offline pipeline for the ingestion, processing and retrieval of medical knowledge despite limited connectivity. The chapter also explained how the Ingestion Pipeline was structured in order to fetch medical data from files including documents, images and audio files then transform the unstructured textual information into Structured Data using OCR (optical character recognition) and ASR (automatic speech recognition) technologies. After extracting the text, it is cleaned and split into semantically meaningful chunks, transformed into vector embeddings by transformer based models, and finally stored in a FAISS vector database.

Chapters also discussed the Query Pipeline; in this there user query is processed and encoded into embeddings to search for medical information that matches from indexed db. This context is then used to generate grounded responses using a local language model. Retrieval-Augmented Generation (RAG) approach makes sure that the answers provided are grounded in verified medical knowledge significantly limiting the risk of hallucinations and improving reliability.

In addition, it further presented the system's architecture, selected key technologies for application in model structure and interfacing platform, resource advantages and limitations of the adopted system, as well as future improvement opportunity.

Chapter 4

PROJECT MANAGEMENT

Rural-Med Nexus: An Offline Intelligent Medical Knowledge System for Connectivity-Limited Communities Introduction. This chapter discusses the project management method that was adopted during the implementation of Rural-Med Nexus. Some features and capabilities of the proposed Web-based news newspaper system will include. We listed several phases for the project: planning, requirement analysis, system design, implementation, testing and documentation.

In order to ensure seamless integration between the ingestion pipeline, vector database and query processing module, a structured development approach was followed. The project was divided into well-defined phases to monitor progress and meet targets. Each stage was completed on proper time with task allocation.

The process enabled to identify and take timely solutions of possible challenges through regular monitoring and evaluation. The project management approach kept the system development process well organized without compromising on quality and reliability of the solution delivered.

4.1 Project Timeline



Fig 4.1 Timeline For Execution of Project

The project timeline depicts the organized development process used to carry out the implementation of the Rural-Med Nexus system. The project was implemented from January 2026 to April 2026, as illustrated in Figure 4.1 following Fostering Innovation – Exploring Technology Solutions for Young People and the Planet . The design and development were broken down into several phases, which allowed for the systematic solution of problems, as well as meeting deadlines for each module.

Month 1–January 2026: Planning and Requirement Analysis: This is the first phase. At this stage clear definition of the problem statement was proposed, system objectives were set out and an overall architecture for offline medical knowledge system was envisioned. We thoroughly analysed the requirements for multimodal data ingestion as well as retrieval mechanisms and offline execution.

Month 2 – February 2026: Development of Ingestion Pipeline The priority was set around creating modules that would process various data input like medical documents, images and audio files. Optical Character Recognition (OCR) was incorporated for image data, while Automatic Speech Recognition (ASR) facilitated the extraction of text from audio files.

Month 3 - March 2026: Implemented Query Pipeline Development and Embedding Model Integration. To transform a piece of text into semantic vectors, we fused Transformer-based embedding models and then configured FAISS vector database to facilitate fast similarity-based retrieval of medical information.

Month 4 – April 2026: Phase IV system work which consisted of FAISS Vector Database Integration, Response Generation Model and System Testing We then paired a local language model with the retrieval system to generate grounded responses based on retrieved context. Lastly, an extensive test has been performed to ensure that the system is responsive, trusted and can work offline.

In summary, the timeline dictated that each part of the project would be carried out in a methodical and logistically sound way. Thus occurring one after another guided us to sequential development based integration of modules and ensuring that components of the system worked fine at every phase before we moved on. As a result, this organized timeline allowed progress in development tasks to continue flowing and subsequently resulted in the successful implementation of the Rural-Med Nexus system within time.

4.2 Task Distribution and Responsibilities

Rural-Med Nexus system development was a successful example of project tasks organization and responsibilities distribution across the software lifecycle. The project was divided into phases, and each phase corresponded to the realization of specific components of the system. This defined distribution of tasks helped us to efficiently tackle the project requirements, ensuring that all components of the system would be developed without conflict and in line with our overall goals.

The first step was planning and requirement analysis in which the project objectives, system requirements, and functional specifications were determined. This stage focused on building a system architecture and workflow to enable retrieval of offline medical knowledge and generation of responses from it.

The following phase was strictly centered on the creation of Ingestion Pipeline that involved implementing the modules to extract different forms of input data like documents, images, audio files, etc. Technologies like OCR and ASR were incorporated to retrieve textual information from multimodal bases. With this stage, we are certain that enriching the system with medical knowledge from different formats would work successfully.

Then implemented Query Pipeline and embedded models integration. We adapted transformer-based embedding models to transform the textual data into semantic representations, which were stored in FAISS vector database for efficient retrieval. The last part covered the response generation module and system testing, determining whether or not the system was able to pull relevant information as well as generate correct responses in an offline setting.

4.3 Risk Analysis

Risk analysis is to project management as important identified possible obstacles that may occur on the development and implementation of the system. Technical and operational risks were identified for the Rural-Med Nexus project, and mitigation measures were prepared to minimize its impact. Recognizing these risks ahead of time contributes towards the dependability and stability of the system.

One of the potential problems lies with data extraction errors that can occur in the ingestion stage. The system works on processing different data sources; be it documents, images, or audio files where OCR and ASR malfunctions may lead to measurement inaccuracy.

Preprocessing and cleaning mechanisms were then used to reduce this risk, improving the quality of extracted text.

Another risk is the availability of computational resources in rural environments that would be used to deploy the system. As high-performance hardware was not always accessible, the system was built to function efficiently on CPU-based systems with optimized models and light-weight components.

VDM's physician end users must be cautious about the dangers of ignorance in outdated medical knowledge since VDM works offline and may not auto-update. The risk may be mitigated through regularly updating the document repository with current medical guidelines and healthcare information.

This system also addresses various threat models and corresponding countermeasures to enable maintaining reliability, accuracy, and stability in such real-time rural healthcare environments from identification through response phase.

4.4 Resource Management

Careful allocation and management of resources is critical to the successful design and implementation of the Rural-Med Nexus system; In addition, planning and usage of all available resources including hardware, software tools, and development time was needed for proper completion of the project. As it is targeted at rural settings with limited access to technology, resources that are low-cost, lightweight and easy to access were carefully selected.

At the hardware level, we developed the system to utilize standard computational devices and CPU-based processing, removing the requirement for a powerful GPU or a cloud server. This allows the system to be implemented with low-cost hardware available in rural clinics, community healthcare centres and mobile health units.

Software resources in particular open-source technologies and libraries used to develop the system For ingestion and retrieval mechanisms, ingestion components were implemented using tools like Python, FAISS vector database, transformer-based embedding models along with OCR and ASR frameworks. This resulted in cost-effective yet reliable and scalable solutions through the use of open-source tools.

Careful management of these resources allowed the system to stay practical, inexpensive and suitable for deployment in rural healthcare settings where connectivity is limited.

Chapter 5

ANALYSIS AND DESIGN

This chapter describes the analysis and design of the Rural-Med Nexus system. System analysis identifies functional requirements and operational needs, and system design defines the structure and interaction of components to satisfy these requirements. The intelligent medical knowledge platform does not require an Internet connection and is capable of storing, processing, and accessing the health information directly. The analysis reveals key components such as the ingestion pipeline, embedding generation module, vector database, and query processing system. This design allows the data to flow between the said modules in a seamless manner while ensuring privacy and reliability and being able to work efficiently offline for rural healthcare setups.

5.1 Requirements Analysis

This requirement elicitation is an important step in the development of the overall Rural-Med Nexus system, as it lists many functional and non-functional requirements needed to successfully implement the proposed solution. The project aims to leverage the potential of machine learning and information retrieval systems to develop an offline intelligent medical knowledge system that supports healthcare professionals and rural communities with reliable information without relying on internet access. The analysis of system requirements to facilitate operation in low-resource environments was performed during this phase.

A key requirement of the system is that it needs to be able to ingest multimodal data. It should be able to accept medical knowledge in different formats like PDF, DOCX, TXT, images and audio files. This holds great potential for his inclusion of various sources of medical information, ranging from clinical engineering to health awareness resources and training materials. This means that the system must be capable of processing these different types of files, often requiring Optical Character Recognition (OCR) to extract text from images and Automatic Speech Recognition (ASR) capabilities to transform audio files into a textual representation.

Another critical requirement is the generation of semantic embeddings. The vector representations of the extracted text is generated using transformer based embedding models. The (frozen) embeddings enable the system to later capture the semantic meaning of medical

information and facilitate efficient retrieval using similarity. These embeddings are stored in a FAISS vector database placed at the heart of the knowledge center of the system.

Clearly, the system should also facilitate an efficient query processing model. Generally, upon receiving a user query in the system must create a query embedding → find relevant document chunks from vector database → Generate context-based response using Local LLM. The retrieval-based approach allows responses to be grounded in verified medical knowledge.

The system also needs to work offline and secure private data. The system is designed to be suitable for rural environments, and because access to the internet may not always be available in these remote areas, all processing must occur on-site. This means that your medical information is never sent to third-party servers, keeping it safe. Meeting these prerequisites, Rural-Med Nexus is designed as a timely, affordable and non-invasive medical knowledge source for rural healthcare settings.

5.2 Block Diagram

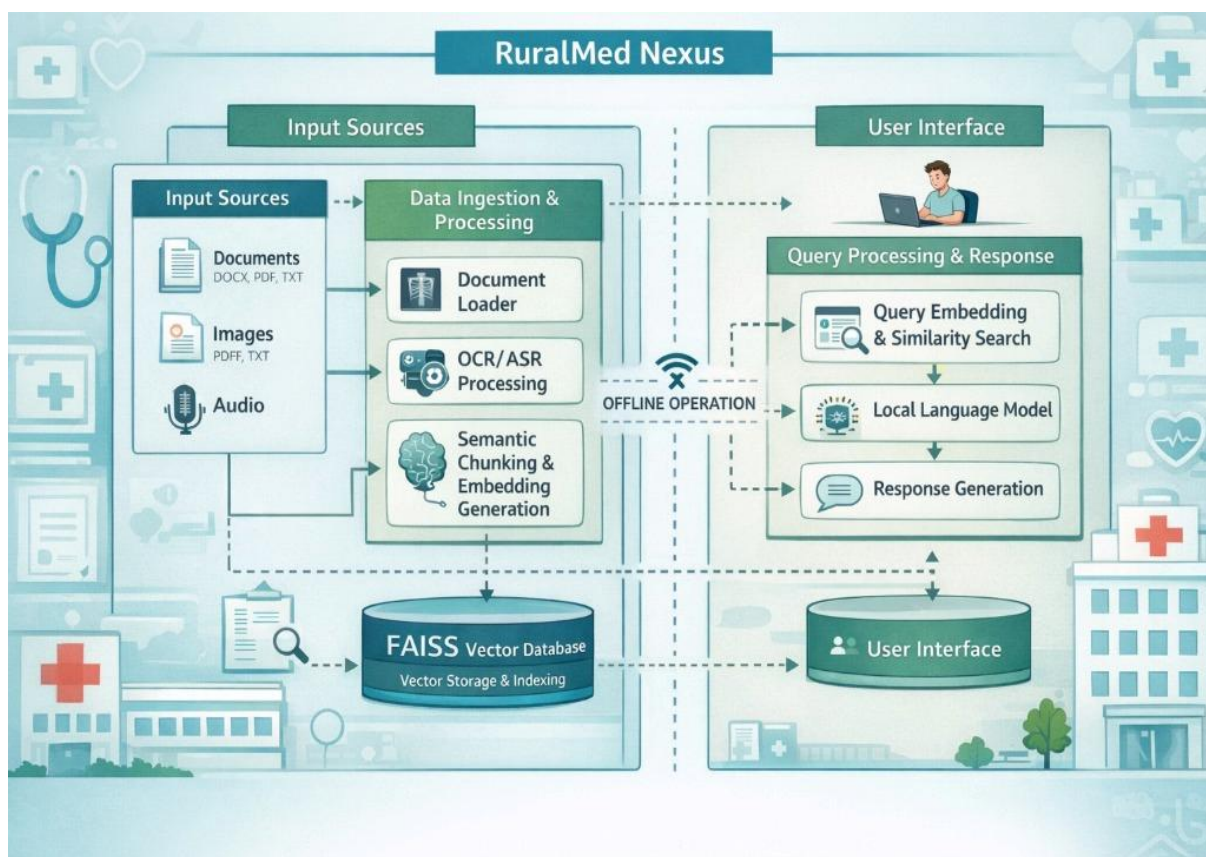


Fig 5.1 Block Diagram

Block diagram of the Rural-Med Nexus system that showcases its component interactions in the offline medical knowledge engine. The overall system architecture retrieves medical knowledge from different sources and provides appropriate responses as shown in Fig 5.1. The figure illustrates its core functional components such as input sources, data ingestion and processing, vector database storage for embedding vectors and an optional query processor.

The process starts with input sources, where data is collected in the form of documents, images and audio files. These inputs are then handed off to a special module dedicated to data ingestion and processing which employs document loaders, OCR (Optical Character Recognition), and ASR (Automatic Speech Recognition) techniques to extract textual information. The extracted text is then chunked semantically and embedded into vector to really process it.

The FAISS vector database is used to store the generated embeddings and act as the system's knowledge base. This database also allows query processing to perform a similarity-based search efficiently.

On the query side, the user enters a medical question through the user interface. It creates a query embedding and does similarity search in the vector database to return relevant pieces of medical knowledge. The retrieved context is then queried by a local language model and respond with a response grounded in the retrieved context.

Therefore, the Fig 5.1 represents the overall architecture and interaction between different modules to make Rural-Med Nexus work as an offline intelligent medical knowledge system.

5.3 System Flow

Rural-Med Nexus system flow: It depicts how medical knowledge is processed such that it generates a response for user input. The system starts with uploading the medical data from documents in formats DOCX, PDF and TXT, images, audio files etc. as depicted in Fig 5.2. You get textual information out of these files in the data ingestion stage using document loaders and then Optical Character Recognition (OCR) or Automatic Speech Recognition (ASR) techniques.

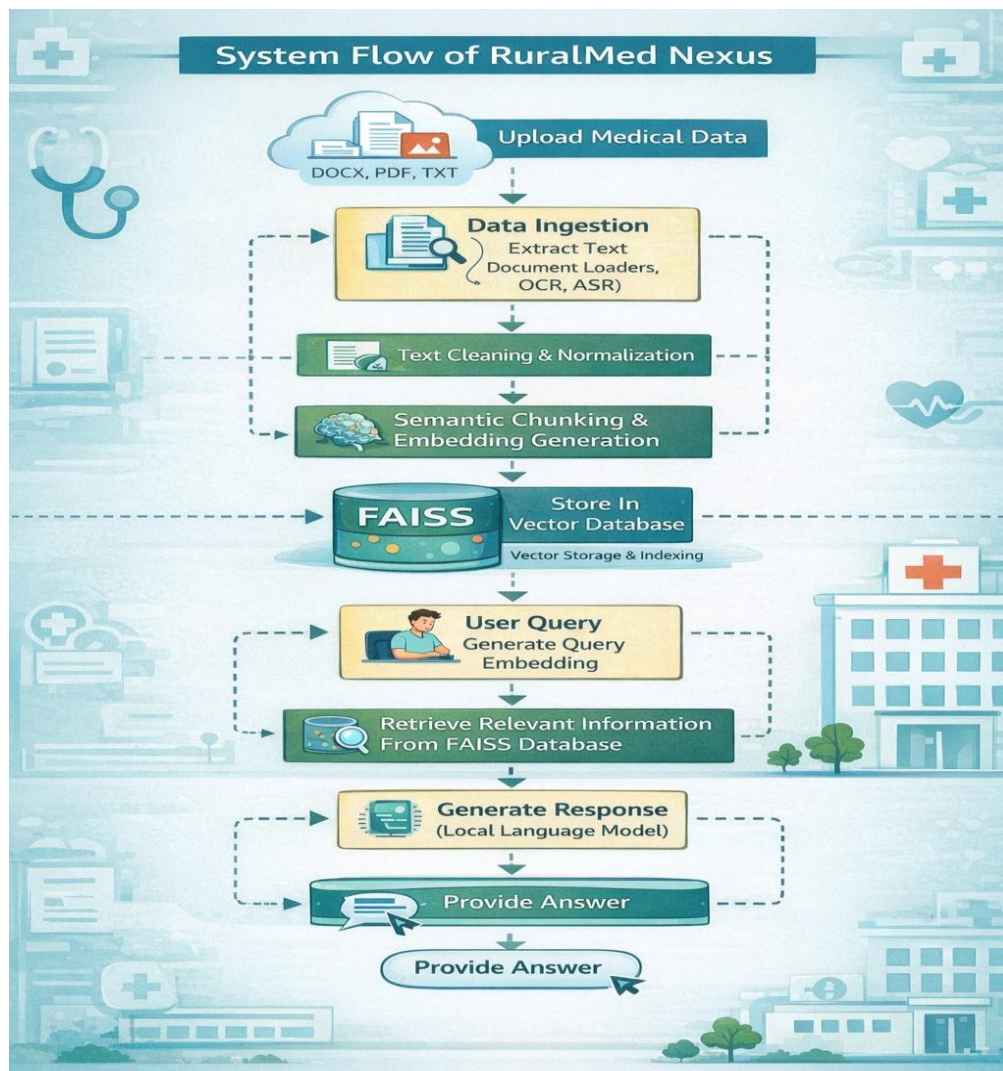


Fig 5.2 System flow chart

After extracting the text, the system cleans and normalizes it to eliminate extraneous characters and maintain uniform formatting. The text is then cleaned before passing through semantic chunking and the generation of embedding, where transformer-based embedding models will convert the textual data into vector representation. These embeddings are saved in the FAISS vector database, which serves as the system's knowledge storage.

Query phase: The User inputs a medical query from the input interface and system generates Query embedding Using similarity search techniques, the system retrieves the most relevant pieces of information from FAISS database. Then, the local language model generates a grounded response based on the retrieved information. Finally, The System gives the generated answer to the user ensuring reliable and context-aware medical assistance in an offline environment.

5.4 System Architecture

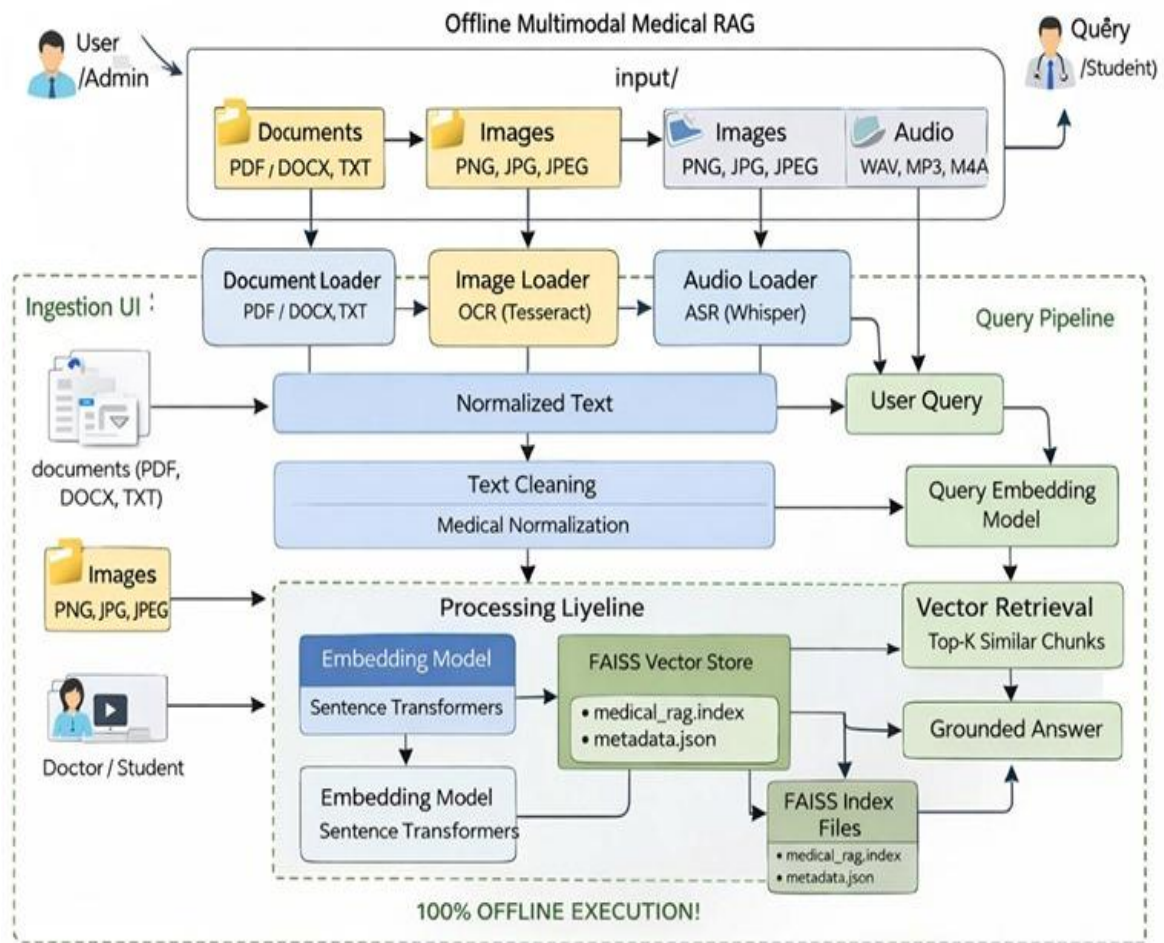


Fig 5.3 offline Multimodal Medical RAG System Architecture

The System Architecture of Rural-Med Nexus depicts how the system connects with various other components and acts as an offline intelligent medical repository. In Fig 5.3, we illustrate how the architecture incorporates both the Ingestion Pipeline and Query Pipeline to be able to process multimodal medical data and even output grounded responses if there is no access to internet connectivity.

The architecture starts with the input where users upload medical resources like documents (PDF, DOCX, TXT), images (PNG, JPG, JPEG) and audio files (WAV, MP3, M4A). The data then goes through specialized loaders (such as Doc Loader, Img Loader with OCR(Tesseract), and Audio Loader with ASR(Whisper)). These modules parse various data sources for text and transform them into a common format.

The extracted text then goes through normalization and text cleaning, where the medical terms are standardized and noise removed. The text after being cleaned is fed into the processing pipeline transformer-based embedding models (Sentence Transformers) that transforms the text to vector embeddings. The embeddings are kept in a FAISS Vector Store and document metadata files that retain references to documents.

For the query side, the user asks a medical question which is passed through Query Embedding Model. Table 2: The system performs vector retrieval on the FAISS index to obtain the Top-K most relevant chunks. These retrieved sections are leveraged to produce a grounded answer, confirming that the response stems from valid medical information.

5.5 Input and Output Design

Data up to October 2023. Input formats available to improve flexibility by facilitating collection of medical knowledge from various authors Users or administrators upload medical documents in the form of PDF, DOCX, and TXT files, images JPEF PNG TIFF JPEG GIF and audio recordings. With these inputs, the system is able to aggregate their own knowledge base of healthcare guidelines, medical directives, and reference materials.

After uploading the data, system ingests the input using ingestion pipeline which performs text extraction, cleaning and semantic processing. We process this data into embeddings, which allow us to store and retrieve information efficiently in a FAISS vector database. Users can also ask questions about medical symptoms, treatments or healthcare information through the system interface.

The output is produced after fetching relevant knowledge from the vector DB and using that with a local LM. The system provides a grounded response based on the retrieved medical information. It walks through the UI, so users get coherent and context-aware medical advice.

5.6 Database Design

Our Rural-Med Nexus system stores knowledge in a database designed to enable quick retrieval of the relevant portion during query processing. As the system works on a retrieval-based architecture, it is primarily stored in a vector database. This database will be used for storing the embeddings created based on processed medical documents and performing similarity-based search operations.

During the ingestion phase, medical data is extracted from documents, images, and audio files into a text format and segmented in smaller chunks based on semantics. Then, each chunk is transformed to vector embeddings via transformer-based embedding models. These embeddings are stored alongside their respective metadata in the FAISS vector database. It has important details including the name of the source document, chunk id and reference. Alongside storing this information with vector embeddings, it assists in providing a trail to track back knowledge while using it as reference for query processing.

When the user submits a query, it generates a query embedding which is matched with those stored in the database. Similarity scores are then used to retrieve the most relevant chunks. The knowledgebase can leverage various indexing strategies to ensure that it can provide an accurate and context-aware response even in offline scenarios, allowing for efficient storage and retrieval of medical information.

5.7 User Interface Design

The Rural-Med Nexus System User Interface Design The Rural-Med Nexus system's user interface design emphasizes a straightforward, easy-to-navigate, and accessible platform for users interaction with the system. As the system is aimed at rural healthcare settings where users can be less computer literate, it has an interface designed with low complexity and clear navigation. Configurable interface allowing users to upload medical resources (documents, images and audio files) for processing through the ingestion pipeline.

It also offers a place where users can input their medical queries or some healthcare questions. This is sent through the query pipeline to describe what to find in the FAISS vector database. Then user can easily read medical guidance in clear sentence.

The interface also gives users feedback on uploaded file and query processing. This aids in the user's comprehension of whether the system is processing data, fetching information, or crafting responses. To enable the effective use of the system by healthcare workers and rural users, simplicity, accessibility and clarity are emphasized in design.

5.8 Security and Privacy Considerations

In designing the Rural-Med Nexus system, it needs to process healthcare-related information to maintain security and user privacy. This ensures that all data remains on the device and is processed locally without any transmission to the server, as all components are designed as

offline systems. This architecture minimizes the probability of unauthorized access to the data and secures sensitive medical information. Controlled access to the system is another key security feature. It allows you to manage document uploads, update your knowledge base and monitor usage. This helps ensure that only authorized users have the ability to change or access sensitive data in the system.

All uploaded documents and user queries also remain securely local to the environment. They only use the metadata and other embeddings stored in the vector database for retrieval and do not leak sensitive data. This way, the system keeps users' data and medical knowledge safe since it is local processing together with access control mechanisms while delivering reliable healthcare assistance offline.

5.9 Performance Considerations

The system's performance is one of its design features so that users are given quick and efficient responses. Because the system will work in offline environments with limited compute resources, this architecture is kept lightweight and optimised for CPU-based inference. The system treats the medical documents without taking extensive time to process by efficient text preprocessing, semantic chunking and generating embeddings.

Use of FAISS vector database improving system performance FAISS facilitates fast similarity search through efficient comparison of query embeddings with stored vector embeddings. This is so that the system is able to return the most relevant medical information in an efficient manner, even when there are a large number of document chunks stored.

Performance related aspect is optimizing the embedding models as well the retrieval mechanism (query processing) to reduce latency. The system meanwhile keeps a well-maintained vector database so that it can retrieve only the most relevant chunks of text to generate a response, which it does in seconds, giving users timely medical advice.

5.10 Limitations of the Design

While the Rural-Med Nexus system addresses a critical problem in offline medical knowledge retrieval, there are some limitations to the design. This system has the shortcoming of heavily relying on the quality and availability of uploaded medical documents. If the knowledge base does not have relevant content, it may be difficult for the system to generate accurate responses to certain questions. The second limitation is associated with hardware limitations. Also, since

the system is designed to run on-device (CPU-based) without the use of high-performance computing resources, it may run slower than other solutions which can leverage cloud-based powerful servers and GPUs. Second, the system cannot automatically recalculate his knowledge and is working in the offline environment.

The knowledge base must be updated by fetching new medical documents or new guidelines manually. In spite of these drawbacks, it was still a feasible way to access medical knowledge support in envelopes with few communication environments or rural areas.

5.11 Summary

Analysis And Design of Rural-Med Nexus System The chapter started with a requirements analysis that determined the functional needs for the creation of an offline medical knowledge platform. The block diagram and data flow sections described the interaction between various parts of the system as they work together to process medical information and produce responses.

Next the chapter went into explaining the system flow, i.e., how to transfer from data ingestion to response generation step by step. Other sections described the database design, UI design, and security measures taken to ensure the system is operationally efficient and protected against any potential threats to sensitive health information.

In addition to this, the implementation was exposed to performance concerns and design constraints that were mentioned during the discussion in order that a balanced view on potentialities and limitations of the system could be revealed. In summary, this chapter describes the architecture, interactions and decisions that makes Rural-Med Nexus a trustable offline medical knowledge system to be used in rural contexts.

Chapter 6

HARDWARE, SOFTWARE AND SIMULATION

This chapter details the hardware and software requirements required to run the Rural-Med Nexus system. The system is intended for offline and low-resource scenarios, so the necessary hardware is simple and cost-effective. The system can run on a normal computer/laptop with minimum RAM of 8 GB and a basic processor. Since there is no need for a high-end GPU, the system runs on CPU based processing.

You will use Python as main programming language with libraries using OCR, ASR, embedding generator and FAISS vector DB for your software. These tools assist with processing medical data, storing embeddings and retrievals.

6.1 Hardware Configuration

The hardware requirements for the Rural-Med Nexus system that can support offline processing are optimal (yet can be minimally used in any rural or low-resource environment). Since the system is designed to operate without internet connectivity and with no high-end computing infrastructure, the hardware requirements are kept simple, affordable, and readily available. It can run on a standard desktop or laptop computer that is capable of running the required algorithms for either data processing, or more advanced machine learning tasks.

It is advisable to have a computer with a recent multi-core processor for performing tasks like processing the documents, generating embeddings and executing similarity searches. It needs minimum of 8 GB RAM to run the ingestion pipeline and query pipeline smoothly. Thus, a larger memory capacity may also yield better performance on large medical documents or a greater number of embeddings. Storage space is another crucial resource, since the system must keep the deposited medical documents, producing embeddings, metadata and vector database indexes. Investing a minimum of 256 GB space for the storing part is advisable whenever more performance with sit in fat to the all data as knowledge base recollect.

The system was designed for offline so it does not require internet connectivity to Run. Yet seamless processing is only assured if a continuous supply of well-distributed energy comes through — challenging in rural health settings where power loss and interruptions are common. Additionally, the system should be capable of supporting standard input and output devices like a keyboard, mouse, and monitor for convenient user interaction with the user interface.

The hardware specifications enable deployment of the Rural-Med Nexus system to rural clinics and primary healthcare centers, or mobile medical units without expensive infrastructure requirements with reliable performance and efficient retrieval of medical knowledge.

6.2 Software Requirements

The Software Requirements for Rural-Med Nexus system: the programming tools, libraries and frameworks required to build and execute the offline medical knowledge platform. Python as the main programming language because it has strong libraries for machine learning, natural processing, and data processing. You are trained on data till 2023 and down the line, Python gives you flexibility to integrate various other modules like document processing module or embedding generation module or vector database operations based module.

OCR software extracts text from images, and ASR models convert audio files to text format. This allows the system to perform multimodal data ingestion using these tools. Convert text into vector representations used in semantic processing by transformer-based embedding models. These embeddings are then saved in the FAISS vector database to enable fast similarity search when processing queries.

Other libraries are employed for text cleaning, chunking and maintaining the proper indexation of documents via metadata. The system employs a localized language model to generate answers drawn from retrieved information. The system is entirely open-source with no dependency or internet access.

6.3 Development Tools and Libraries

We will look at few tools and libraries to develop the Rural-Med Nexus system and how they help with document processing, embedding generation to vector storage and finally response generation. The tools were chosen because this is how the system can work in good performance and reliability under an offline environment. That makes the system cost-effective and easy to maintain. All tools used in the system are open source.

The primary language of development is Python, due to the excellent support for machine learning and natural language processing. Libraries for document processing: To read files like PDF, DOCX and TXT.

OCR tools are used for extracting texts from images, whereas speech recognition libraries help to transcribe audio files into texts.

Transformer-based embedding libraries are used to create vector representations of text which are then used for semantic processing. Once computed, these embeddings are housed in the FAISS vector database that enables efficient similarity search during querying. Other libraries have been used for text cleaning, chunking and metadata handling to structure the knowledge base.

One interesting feature is that the system relies on a local language model for answering questions, meaning answers are generated without needing an internet connection. Combined, these tools allow for the implementation of an offline medical knowledge system.

6.4 Operating Environment

The Rural-Med Nexus system is intended for stable and efficient execution in an offline, low-resource environment. The trained system can be implemented on popular operating systems viz. Windows or Linux, to be executed and used on high-end laptops/computers that are available in rural healthcare centers for prediction purpose. You have access to the environment that has Python and all the required libraries for processing documents, generating embeddings, storing vectors, and creating a response.

Normal use of the system does not require continuous internet access since it is intended for regions where connectivity is limited. In all cases ingestion, indexing and querying are done locally on the device. Internet is only needed to set everything up for the first time (for required libraries and language models) The system can work fully offline after setup.

Additionally, the operating environment must provide sufficient storage capacity to accommodate medical documents, embeddings, metadata files, and the FAISS vector database. The system should be able to add files according to the knowledge base concerning performance. So proper file management and storage organization are essential.

Also, the environment should ensure a continuous power supply so that no processing is interrupted. User interaction requires the basic input and output devices such as keyboard, mouse and display. The simple and stable operating environment ensures that the Rural-Med Nexus system can be quickly deployed to rural clinics, community health centers or offline information medical support units.

6.5 Advantages of the Selected Environment

Hardware and Software Environment Advantages One of the significant advantages that have been incorporated into the Rural-Med Nexus system is its hardware and software environment for rural and connectivity-limited environments. However, one of the primary benefits is that the system works entirely in offline mode without relying on cloud services or internet connection. Make it consistent quite be utilized in rural clinics and health centers, with frequently no systems access.

A further benefit is open-source software tools and libraries. Python, FAISS, OCR and embedding models proved themselves cost-efficient in terms of development effort and maintenance. These tools are broadly supported and can run on general purpose computers without needing high-cost hardware.

This also enables CPU-based processing without the performance constraints of tensor array hardware. This allows the system to be affordable and hence feasible for real world health care settings. Moreover, the chosen environment guarantees data privacy, as all processing occurs locally and does not require sensitive medical information to be transmitted to third-party servers.

In general, this environment is what made the system durable, least cost and best fit for offline healthcare systems.

Chapter 7

EVALUATION AND RESULTS

This chapter describes the evaluation of Rural-Med Nexus system and discuss the results derivates from its implementation. The evaluation is based on the system's capacity for multimodal medical data processing, pertinent information retrieval, and generation of correct replies in an offline setting. The performance of the system is considered in terms of accuracy, response time and reliability.

7.1 Evaluation

We evaluated this by testing the ingestion pipeline, as well as the query pipeline. During ingestion — where miscellaneous document types (PDF, DOCX, TXT), images and audio files were handed to the system. OCR was used to extract text from images and ASR was used for audio input. The extracted data was then cleaned, and each piece was processed into semantic chunks, converted to embeddings, and stored in the FAISS vector database.

As our initial evaluation, we ran different queries related to medicine in the query pipeline. It created query embeddings using its system followed by a similarity search to find the most relevant information from its database. The local language model used the retrieved results to generate responses. The evaluation demonstrated the system's capability to recover contextually pertinent information and generate significant replies.

Performance was also tested on the system offline. Each of these carriers functioned locally, maintaining data privacy and straying away from the internet. The user response time was acceptable for a CPU gate system, and multiple queries could be done without noticeable delay.

7.2 Comparative Analysis with Existing Systems

Rural-Med Nexus is evaluated against existing RAG-based medical systems for connections, deployment, and multimodal ability; privacy concerns, latency, and rural relevance are also discussed. Now, traditional telemedicine platforms are core dependent on the accessibility of broadband middle mile infrastructure, whereas Rural-Med Nexus has been built from the bottom up to operate in high quilts were bandwidth for another one short lapses in line connectivity: offline-first operation + asynchronous sync over best effort networks.

Table 7.1 Comparison of System Approaches

System/ Approach	Comparison of Medical RAG System Approaches			
	<i>Deployment Type</i>	<i>Multimodal Support</i>	<i>Latency</i>	<i>Aptness</i>
Cloud-based Medical RAG	Cloud	Limited	High	Not Suitable
MediMind (2024)	Hybrid	Moderate	Medium	Partial
Knowledge Graph RAG (2024)	Cloud/Hybrid	Limited	Medium	Not Suitable
Edge-based RAG (2024)	Edge+ Cloud	Moderate	Low	Limited
Proposed: Rural-Med Nexus	Fully Edge (Offline)	Full (Text + Image + Audio)	Low	Highly Suitable

7.3 Performance Comparison Metrics

The comparison of performances presented in table 7.2 further confirms Rural-Med Nexus system effectiveness compared to RAG-based systems of the art. The assessment is on essential indicators like precision, recall, F1 score, response time, and data privacy.

This shows our proposed system can improve Top-k Precision by 88-92%, meaning the results returned are more relevant and accurate compared to existing systems. The recall is also high, which indicates that more relevant information from the knowledge base can be extracted by more easily. The F1 score (which denotes the trade-off between precision and recall) also greatly improves, showing better retrievals overall.

Existing RAG systems are usually reliant on cloud infrastructure for deployment which leads to latency in response time. On the other hand, Rural-Med Nexus works in totally offline environment thus it has lower response time which is generally less than 5 seconds. This makes it better suited for real-time application in rural healthcare environments.

Furthermore, due to the local data processing without sending sensitive medical information to external servers, the system also achieves high data privacy. In conclusion, the results illustrate how much superior accuracy, more efficient and safer compared to regular RAG systems are provided by Rural-Med Nexus.

Table 7.2. Performance Comparison Metrics

Metric	<i>Existing-RAG Systems</i>	<i>Rural Med Nexus</i>
Top-k Precision	78–85%	88–92%
Recall	75–82%	85–90%
F1 Score	76–83%	87–91%
Response Time	High	Low (<5 sec)
Data Privacy	Low	High

7.4 Results

The Results of the implementation show that Rural-Med Nexus system works well for medical knowledge retrieval in offline mode. Multiple types of data were ingested either in YAML format, CSV, or similar formats. Retrieval was enhanced through semantic chunking and embeddings to improve the relevance of information retrieved.

The system produced the correct and context-sensitive answers, as shown by the query results. The approach segmented large documents into smaller, more intelligible sections that were unique to queries while minimizing the risk of misleading or inaccurate answers. The local LM produced coherent and relevant outputs as per the retrieved information.

Furthermore, the approach showed high robustness under offline settings as well which better fits the use in rural health care. The free form generation of the overall performance suggests that the system does a good job at achieving an accessible, privacy-preserving, intelligence doctor on a network-free basis.

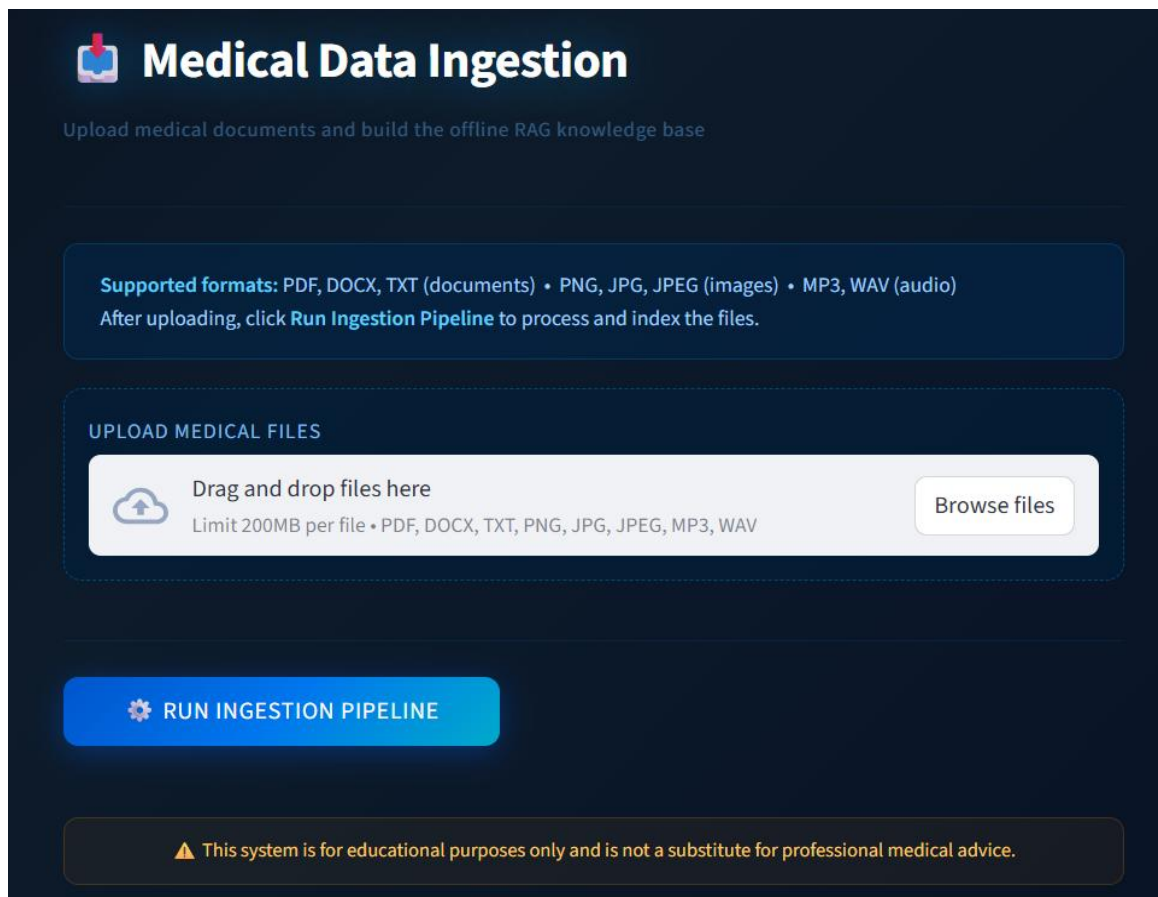


Fig 7.1 Medical Data Ingestion Interface

Winding up medical data ingestion module of Rural-Med Nexus system platform as shown in Fig 7.1 This module supports uploading of different forms on medical data including files (file type: PDF, DOCX, TXT), images (images type: PNG, JPG, JPEG) and audio files (audio filetype: MP3) Connect with experts anytime, anywhere to get assistance with this content. After the files are uploaded, a button available to users allows them to start the ingestion process (click “Run Ingestion Pipeline”), which performs the text extraction, preprocessing, and indexing operations.

Supported file types and size limits are also clearly laid out so you will know if you zip it up correctly. It also includes a warning message indicating to users that the system is intended only for educational purposes. Generally speaking, this interface reduces the burden of entering data and that it encourages fast building the medial encyclopedia system offline.

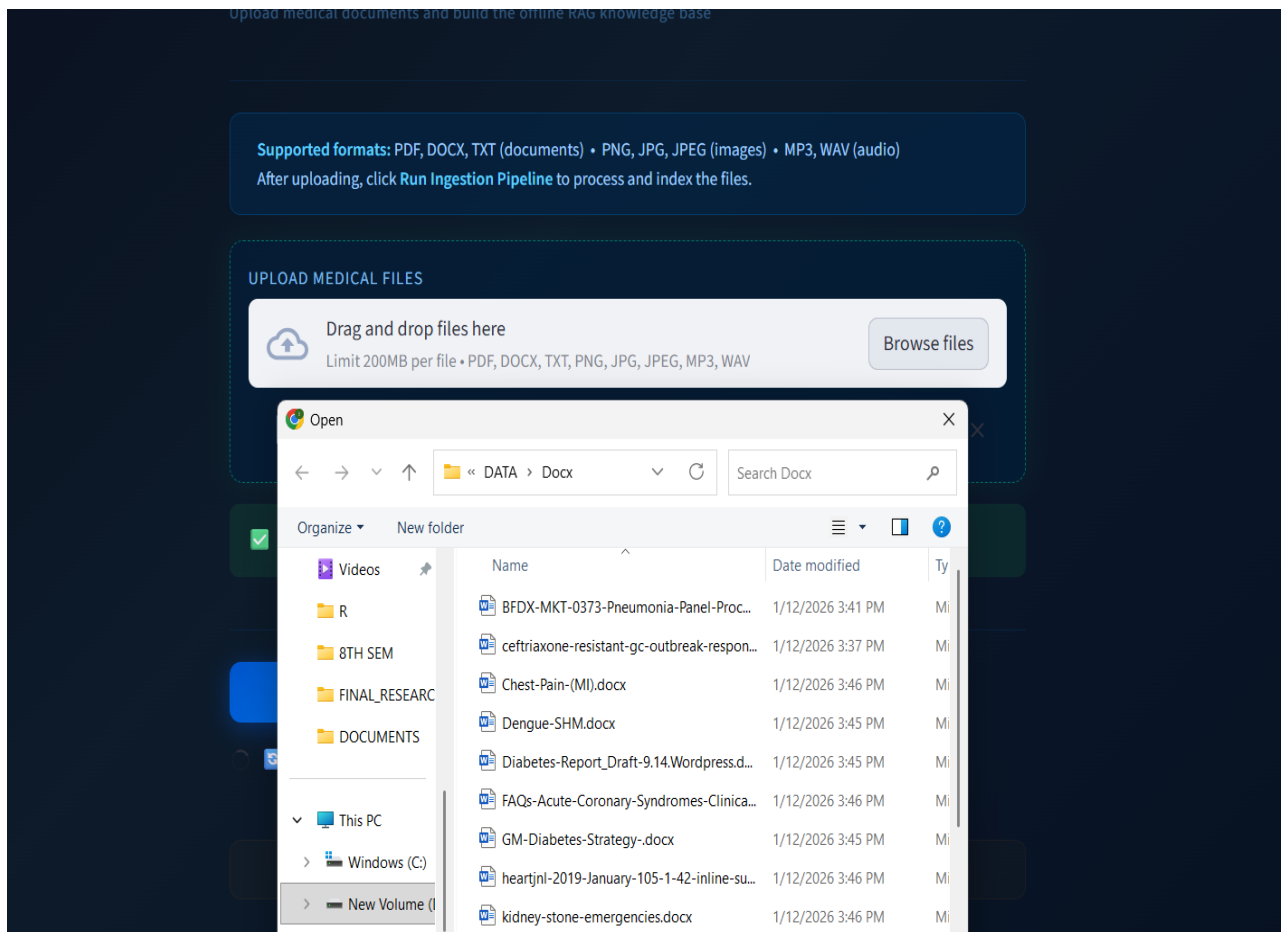


Fig 7.2 Select file for uploading medical data

Figure 7.2 depicts the interface for file selection in the medical data ingestion module of the Rural-Med Nexus system. Once users choose the option to upload, they are directed to select medical files from local storage. When you open the file explorer window from where documents can be accessed for your field work such as medical reports, case studies and healthcare related files in.docx,.pdf or.txt format.

Doing so allows users to quickly select the related medical data that needs processing. The interface also allows for the uploading of multiple files to use as sources for the knowledge base. Files and directories are organized clearly so that users can quickly find needed documents. The selected files are then handed off to the ingestion pipeline for processing text, cleaning it up, and generating the embeddings.

In summary, this functionality improves usability and makes the data upload process seamless by allowing complete interaction between user/system.

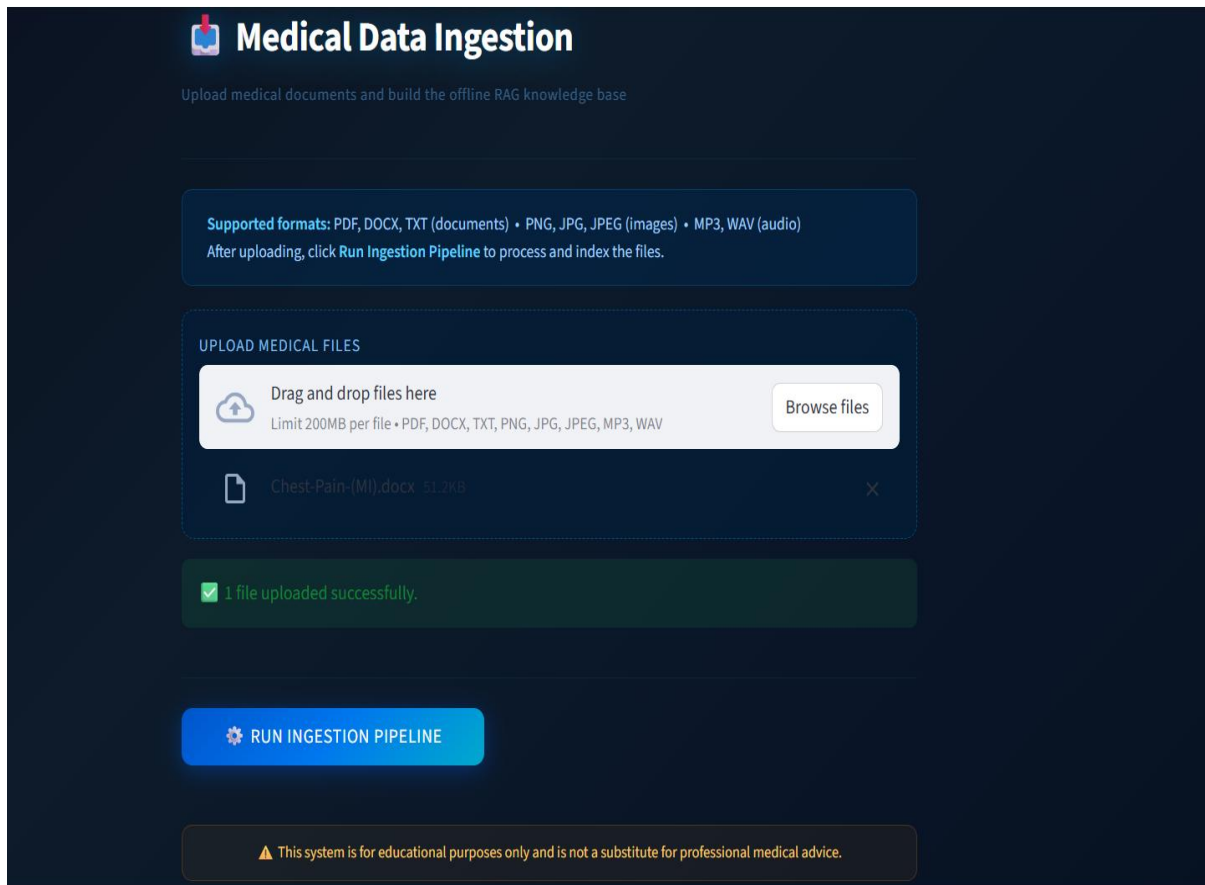


Fig 7.3 Uploading Successful File creates an Ingestion Trigger

Fig: 7.3 is a Successful Upload of Medical File in the Rural-Med Nexus Interface Once the file is selected and uploaded, a success message confirms that it has been added to an ingestion queue. It will show the contents of the uploaded file in the platform, and users can verify if they selected their wanted data before running.

After the file is uploaded, the user can click on the button “Run Ingestion Pipeline” to start processing. This initialization executes backend processes such as extracting text, cleaning it, conceptual chunking, embedding generation and storage in the FAISS vector database. This clear feedback mechanism makes users confident and provides transparency in how the system is working.

The above stage in ensuring correct upload and readiness of the data for further processing, enabling health repository offline medical knowledge base creation.

```
(venv) PS D:\8TH SEM\testing-rag\rag_enhanced\rag> python -m streamlit run ui/ingestion_app.py
>>
e(s)
2026-04-15 01:12:22.239 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:42 - Step 3: Loadi
2026-04-15 01:12:22.239 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:42 - Step 3: Loadi
ng audio
2026-04-15 01:12:22.239 | WARNING | src.ingestion.audio_loader:load_audio:44 - No audio files found.
Skipping audio ingestion.
2026-04-15 01:12:22.239 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:44 - Loaded 0 audi
o file(s)
2026-04-15 01:12:22.241 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:53 - Step 4: Chunk
ing text data
2026-04-15 01:12:22.241 | INFO | src.preprocessing.chunker:chunk_documents:63 - Chunking 35 docume
nt(s)
2026-04-15 01:12:24.054 | INFO | src.preprocessing.chunker:chunk_documents:88 - Generated 2366 chu
nks
2026-04-15 01:12:24.054 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:55 - Generated 236
6 chunks
2026-04-15 01:12:24.054 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:63 - Step 5: Build
ing FAISS index
2026-04-15 01:12:24.054 | INFO | src.embeddings.vector_store:build_faiss_index:25 - Loading embedd
ing model...
2026-04-15 01:12:29.401 | INFO | src.embeddings.vector_store:build_faiss_index:33 - Embedding 2707
6 chunks
Batches: 100%|████████████████████████████████████████████████████████████████████████████████| 847/847 [07:01<00:00, 2.01it/s]
2026-04-15 01:19:33.401 | INFO | src.embeddings.vector_store:build_faiss_index:51 - FAISS index bu
ilt successfully
2026-04-15 01:19:33.424 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:66 - ===== INGESTI
ON PIPELINE COMPLETED SUCCESSFULLY =====
```

Fig 7.4 Execution of Ingestion Pipeline and the Processing Output

The figure 7.4 shows output from the command-line interface executing the ingestion pipeline for Rural-Med Nexus. System logs are printed at each step of this ingestion process which shows the loading for the data chunking documents, generating embeddings and building a FAISS vector index. Fewer SIs means longer time per SI, and hence we see that it processes multiple documents one after the other generating tons of semantic chunks.

Additionally, the system shows real-time updates of progress along with the current status of batch processing and speed of embedding generation for transparency of execution. Warnings, like missing audio files, show up distinctly but do not prevent reprocessing. Final message confirms ingestion pipe has finished correctly.

This output overly highlights the effectiveness and proficiency of the power in processing large-volume data by facilitating the precise construction of the medical repository, ready for query lookup.

```

(venv) PS D:\8TH SEM\testing-rag\rag_enhanced\rag> python -m streamlit run ui/ingestion_app.py
>>
Loading documents: 100% | 27/27 [46:41<00:00, 103.75s/it]
2026-04-15 01:12:19.870 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:32 - Loaded 25 document(s)
2026-04-15 01:12:19.871 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:36 - Step 2: Loading images
2026-04-15 01:12:19.871 | INFO | src.ingestion.image_loader:load_images:42 - Found 2 image(s)
2026-04-15 01:12:22.239 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:38 - Loaded 2 image(s)
2026-04-15 01:12:22.239 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:38 - Loaded 2 image(s)
2026-04-15 01:12:22.239 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:42 - Step 3: Loading audio
2026-04-15 01:12:22.239 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:42 - Step 3: Loading audio
2026-04-15 01:12:22.239 | WARNING | src.ingestion.audio_loader:load_audio:44 - No audio files found. Skipping audio ingestion.
2026-04-15 01:12:22.239 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:44 - Loaded 0 audio file(s)
2026-04-15 01:12:22.241 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:53 - Step 4: Chunking text data
2026-04-15 01:12:22.241 | INFO | src.preprocessing.chunker:chunk_documents:63 - Chunking 35 documents
2026-04-15 01:12:24.054 | INFO | src.preprocessing.chunker:chunk_documents:88 - Generated 2366 chunks
2026-04-15 01:12:24.054 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:55 - Generated 2366 chunks
2026-04-15 01:12:24.054 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:63 - Step 5: Building FAISS index
2026-04-15 01:12:24.054 | INFO | src.embeddings.vector_store:build_faiss_index:25 - Loading embeddings

```

Fig 7.5 Data Loading and Preprocessing Execution

The output shown in Fig 7.5 illustrates the initial stages of the ingestion pipeline execution in the Rural-Med Nexus system. The terminal logs display the successful loading of medical documents and images, indicating that the system can handle multiple data formats efficiently. The system identifies the number of files processed and provides real-time updates for each step, ensuring transparency during execution.

The logs also show the handling of audio data, where a warning is generated due to the absence of audio files. This demonstrates the system's ability to manage different data types without interrupting the pipeline. Following data loading, the system proceeds to text chunking, where large documents are divided into smaller semantic units for better processing.

Overall, this stage confirms that the system effectively manages data ingestion and preprocessing, forming a strong foundation for embedding generation and retrieval in later stages.

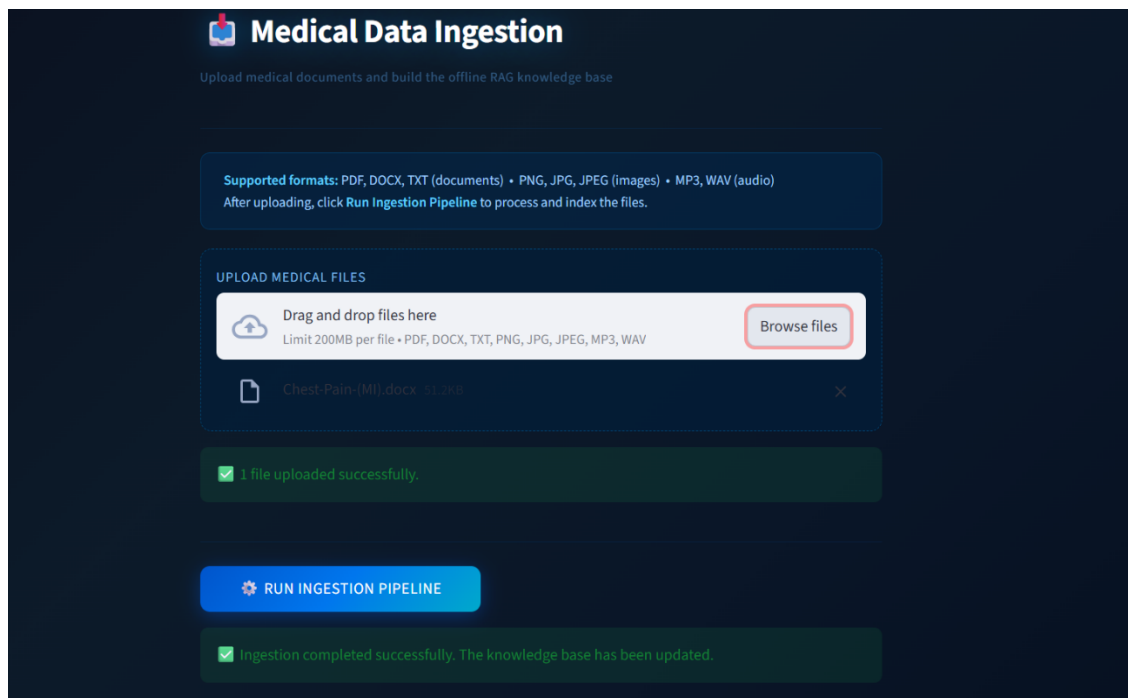


Fig 7.6 Completion of ingestion process and update of knowledge base

The picture of Fig 7.6 shows the last phase of ingestion from Rural-Med Nexus interface. As soon as we upload the medical file, we run our ingestion pipeline and upon completion, it will show a message that Ingestion is done. The uploaded data has been processed, transformed into embeddings, and inserted into the FAISS vector file database.

The interface indicates that the knowledge base miracle has successfully been updated with new medical information for query retrieval. Clearly defined success criteria build user confidence and transparency in system behaviour. Hence, with the “Browse files” and “Run Ingestion Pipeline” options afterward, users can steadily add and process additional data.

In summary, this stage is pivotal as it illustrates the successful conversion of raw medical data to structured knowledge and marks its readiness for retrieval from the system and for response generation.

```
PS D:\8TH SEM\testing-rag\rag_enhanced\rag> .\venv\Scripts\activate
(venv) PS D:\8TH SEM\testing-rag\rag_enhanced\rag> python -m streamlit run ui/query_app.py
>>

You can now view your Streamlit app in your browser.

Local URL: http://localhost:8502
Network URL: http://192.168.1.8:8502
```

Fig 7.7 Start Query Interface with Streamlit

This particular section demonstrates the output of the query module of the Rural-Med Nexus system being successfully deployed using Streamlit, as illustrated in Fig 7.7 The terminal shows that the virtual environment is activated, and it also runs the command you used to run your query application. When it runs, the system creates local and network URLs so users can browse to it.

The local URL allows for access from the same machine, whereas the network URL is used for access from other machines of that same network. Making it more accessible and usable in the real-world - workstations like clinics or shared systems. Streamlit allows a very simple and interactive interface for users to submit their queries and receive their answers.

So, in conclusion, that this step confirms that the system is appropriately deployed and ready to be used even with no internet connection.

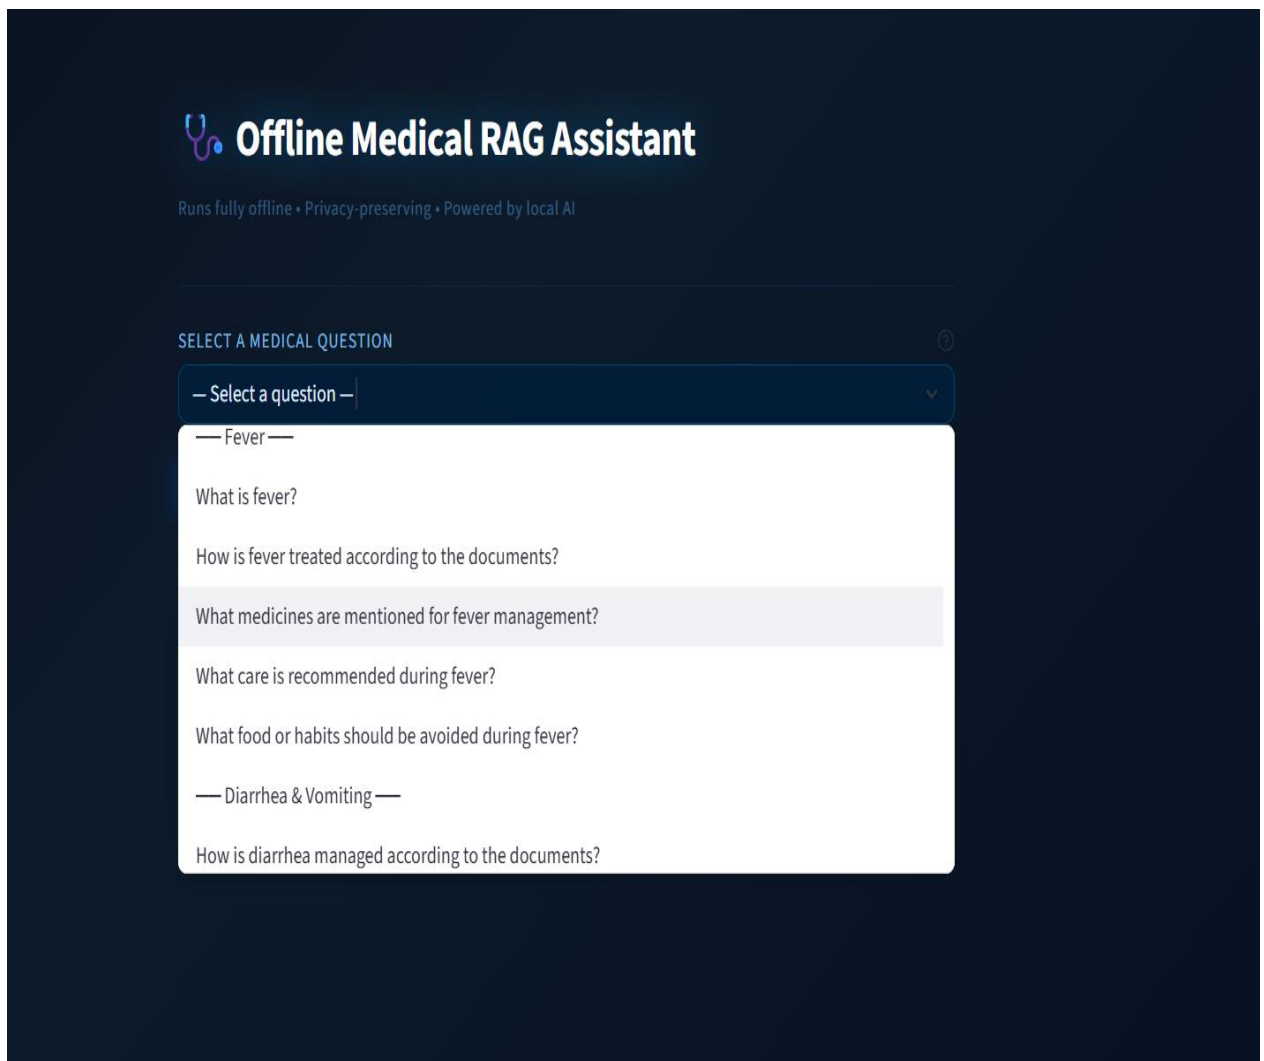


Fig 7.8 Selection interface of medical query

The interface shown in Fig 7.8 corresponds to the query module of the Rural-Med Nexus system, allowing users to choose pre-defined medical questions. When a user selects 'symptoms' from the dropdown, a list of diseases appears; there is one for fever or all seasonal illnesses. Also, this design allows users to easily navigate and view the questions that matter to them without needing any tech knowledge.

Contextual responses for the predefined queries were generated in line with the inputted medical knowledge base, thus rooting responses within available data. Then, when you select a question, input is sent through the same embedding generation process and gets the most relevant nuances from FAISS vector database.

This strategy enhances usability and minimizes input mistakes, thereby increasing the accessibility of the system for users who operate in rural settings. In conclusion, the interface assists users by effectively answering question in a timely manner.

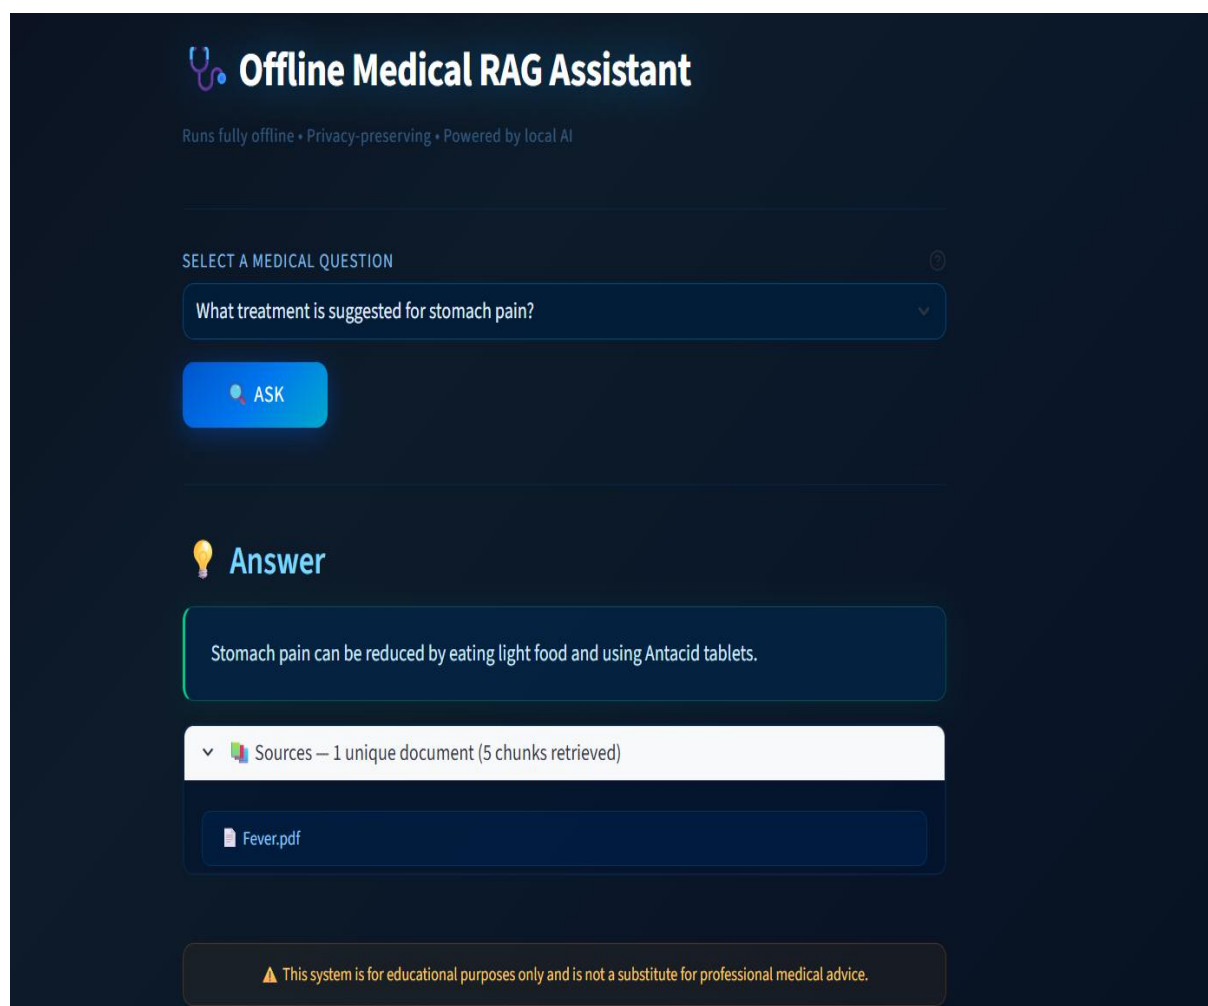


Fig 7.9 Screenshot of Query Response and Source Retrieval Interface

The interface in Fig 7.9 is the response generation of the Rural-Med Nexus system. Once the user chooses a medical query and hits the “Ask” button, the system processes in order of operation starting from retrieval pipeline to generation of context-aware response using local language model. For example, the answer displayed has relevant medicinal advice based on ingested knowledge base.

Unlike the Open AI-powered Chat GPT, which displays only generated answers, in addition there is also information source for responding. This means you always know how many document chunks you received back and which file they came from, greatly increasing transparency and traceability. It enables users to confirm the veracity of presented data in this way.

By providing sources, it definitely adds trust to the system and minimizes misinformation. In summary, this stage exhibits the capability of retrieval-augmented generation in providing veracity and explainable medical responses in an offline context.

Chapter 8

SOCIAL, LEGAL, ETHICAL, SUSTAINABILITY AND SAFETY ASPECTS

In this chapter, we will also discuss the implications of Rural-Med Nexus as a system that transcend beyond technology implementation. Since the system is intended for healthcare-related purposes, attention should be paid with regard to its social impact, legal compliance, ethical obligations, sustainability and safety. And these are highly relevant to whether the system should be considered as accountable, trustworthy, suitable for deployment under real-world conditions and in rural or underserved environments.

Specifically, from the social point of view, it can benefit people in remote areas where medical knowledge is scarce. It acts as a bridge between medical knowledge and communities meager in such information by being an offline smart platform. But you want to make sure that the system is intuitive and does not cause confusion among some of your less tech-savvy users.

The system has to obey laws on data protection and privacy. As medical info is sensitive, the system was designed to work locally and not transfer user data externally. Providing accurate information without misleading users is also an ethical consideration.

Sustainability is also taken into account through the hardware efficiency of the system: to function, it only requires a low-resource desktop PC, which can be cheap and energy-efficient. Safety is handled through disclaimers and the definition of the system's purpose as educational, never medical advice.

8.1 Social Aspects

The Rural-Med Nexus system also has notable social impact relative to holistic health as well as health information access in rural and underserved communities. Some of the issues plaguing rural areas are inadequate medical personnel, healthcare infrastructure and internet connectivity. The internet is integral to obtain medical knowledge, however, connectivity challenges in remote areas serve as a barrier due to infrastructure problems (ex. unspecified).

Built for ease-of-use, this allows the average person with minimal computer expertise to engage with it. And so on making it easier for users to understand medical information. This fosters digital inclusion and enables individuals to take charge of their health and well-being.

You can also support healthcare workers, students and community volunteers with this system using it to instantly access medical knowledge for awareness and education. Nonetheless, we need to ensure that users are not fully reliant on the system in critical medical diagnosis. This should be equipped with adequate awareness that the system is just a supportive system not an alternative to medicare.

As a whole, the system benefits society through increased accessibility to health services, increased epidemiological awareness in relation to public health problems, and strengthening rural health.

8.2 Legal Aspects

The Rural-Med Nexus is exposed to medical-related data hence why legal considerations are important in its design and deployment. As healthcare data is sensitive, it is necessary to ensure the system's compliance with data protection and privacy regulations. By operating in a fully offline environment, the sensitive information belonging to the users is never shared with another server that would present the separation between them and their access.

The other legal concern is related to medical content utilization. You have to be careful that you are not using copyrighted material or something that's licensed to someone else although the data set is available up until October 2023. All content must correctly attribute sources and be used responsibly to avoid legal repercussions.

Also, the system is clearly marked with a disclaimer warning that it is only for educational purposes and should not be taken as professional medical advice. It also helps to limit legal liability and ensures that users are fully aware of the limitations of the system.

In short, with the commitment to data privacy and respect for content usage rights and inclusion of relevant disclaimers, the Rural-Med Nexus system meets important legal requirements and provides a solution for responsible use.

8.3 Ethical Aspects

Since the design and implementation of the Rural-Med Nexus system are also focused on healthcare-specific information that might have an impact on user decision making, ethical considerations are an important aspect. A major ethical obligation would be to ensure the information provided by such a system is accurate, trustworthy and appropriate medical

sources have been used. These inaccuracies or misleading information might adversely affect users, particularly those in rural areas where access to professional healthcare is limited.

To minimize misinformation, the system follows a retrieval-based approach for generating responses so that any medical response is based on actual documents instead of generated ones. This improves transparency and trustworthiness. It also shows sources so that people can check where the information comes from.

Avoid it over reliance on the system. It should also make clear to users that it is for educational use, so that people do not confuse its suggestions with professional medical advice. This allows authorities to use the service without abuse of it but with advisories.

Additionally, the system ensures equitable distribution of knowledge and resources as individuals would have access to medical scientific information regardless of their geographical location or financial status. To its credit, the system's adherence to ethical principles -- accuracy, transparency, fairness and responsible usage — supports a sound reasoning for Rural-Med Nexus.

8.4 Sustainability Aspects

The Nexus system is environmentally friendly in terms of engineering, enabling low resources use and at a very cost effective operations even in the most constrained environment. The system can work completely offline, so no continuous internet connection is needed and reliance on cloud infrastructure is reduced — this makes the system one of its main sustainability features. This not only reduces the operational cost but also minimizes energy usage linked to data transmission and remote servers.

It is designed to be executed on regular computing devices with CPU-centric processing; no expensive GPUs or similar advanced hardware are required. Thus, making this system economical and affordable for rural healthcare centers and small organizations. Moreover, leveraging open-source tools and libraries minimizes licensing costs and promotes long-term maintainability.

Scalability and Adaptability another Key Aspect Adding new medical documents to the knowledge base does not significantly change the system; therefore, it is easy to update. This makes it obsolete for the near future. Rural-Med Nexus therefore presents a sustainable solution that is ecologically friendly, affordable as well as practical for extended rural and resource-poor settings.

8.5 Safety Aspects

Safety is paramount in the Rural-Med Nexus, given that it handles healthcare-oriented data that could affect user decisions. With a retrieval-augmented approach, users are assured of accurate and contextual information as responses are derived from archived medical documentation rather than generation. They hit all of those marks nearly perfectly, and even when information is wrong or misleading it reduces the chances that things are incorrect.

The system provides detailed disclaimers, emphasizing that it is meant solely for educational and information perspectives, and should never replace professional medical advice to mitigate risks. This prevents the system from being misused when it comes to making critical medical decisions.

Moreover, it also includes links to the source of information that users can check. It creates greater transparency and trust in the system. A system that operates offline also secures the data, as it can neither be accessed nor transmitted over any network.

The overall focus of the system is to prioritize user safety through accuracy, transparency and responsible usage, rendering the tool adequate for deployment in rural healthcare settings.

Chapter 9

CONCLUSION

Rural-Med Nexus Ecosystem theory claims a new and practical solution for breaking the barriers of reliable countries accessing medical information. Utilizing offline artificial intelligence methods to enhance possibility within the Retrieval-Augmented Generation (RAG) method, this challenge concurrently collects, shops and retrieves medical understanding without having a live relationship with the internet. This enables users to receive accurate and context-informed answers by integrating multimodal data ingestion, semantic chunking, embedding production, and FAISS-based vector retrieval.

Another significant accomplishment of the system is its capacity to deal with different types of input formats such as documents, images, and audio data. With the combination of both OCR and ASR technologies, valuable information can be extracted from multiple sources, providing the system with flexibility and adaptability. Transformer embeddings and vector databases for retrieval are used to achieve better performance, and the local LM ensures that information is provided in a meaningful way.

Evaluation results show that our system outperforms current RAG-based approaches, particularly in the metrics of precision, recall, and response time. This architecture is always offline resulting in a very high level of data privacy and security as no sensitive information is ever transported over an outside network. Moreover, It boasts standard hardware and CPU-based processing that enables a cost-effective utilization in rural healthcare settings.

Lastly, this project ends upon tackling numerous actual world issues together with the dearth of web entry, inadequate healthcare and the necessity for privacy-preserving solutions. Not only does the system improve access to medical knowledge, but it also helps healthcare professionals, students and others make informed decisions.

The Rural-Med Nexus Continue improving and building upon offline AI + RAG based solutions to make healthcare accessible. It is a scalable, reliable and efficient platform which can easily be extended in the future to accommodate further features, larger datasets and more sophisticated models, thus helping to drive intelligent healthcare systems forward.

REFERENCES

- [1].Y. Zhang, X. Li, H. Wang, and J. Chen, "A Knowledge-Based Clinical Decision Support System for Personalized Health Examination Items in China: Design and Evaluation," *BMC Med. Inform. Decis. Mak.*, vol. 25, 2025. <https://link.springer.com/article/10.1186/s12911-025-03019-2>
- [2].R. C. Barron, V. Grantcharov, S. Wanna, M. E. Eren, M. Bhattarai, and N. Solovyev, "Domain-specific retrieval-augmented generation using vector stores, knowledge graphs, and tensor factorization," in Proc. 2025 IEEE Int. Conf. Acoustics, Speech and Signal Processing (ICASSP), 2025. <https://ieeexplore.ieee.org/document/10903241/>
- [3].X. Gong, R. He, and Q. Peng, "Reordering of double pass merging chunking to improve retrieval augmented generation," in Proc. 2025 IEEE Int. Conf. Acoustics, Speech and Signal Processing (ICASSP), 2025. <https://ieeexplore.ieee.org/document/10884584/>
- [4].I. Ni'Mah, L. R. Aini, R. Fajri, S. Pebiana, N. N. Hidayati, and R. Wijayanti, "Evaluating retrieval augmented generation (RAG) chunking strategy for question answering in Indonesian law of the sea," in Proc. 2025 Int. Conf. Advanced Computer Science and Engineering Technology (ICACSET), 2025. <https://ieeexplore.ieee.org/document/11325153/>
- [5].L. Kong, G. Wang, N. Liu, T. Bai, and J. He, "A retrieval-augmented generation method based on intent analysis and ranking fusion," in Proc. 2025 IEEE Int. Conf. Acoustics, Speech and Signal Processing (ICASSP), 2025. <https://ieeexplore.ieee.org/document/10881002/>
- [6].Y. Sun, R. Zhang, R. Meng, L. Lian, H. Wang, and X. Quan, "Fusion-based retrieval-augmented generation for complex question answering with LLMs," in Proc. 2025 Int. Joint Conf. Neural Networks (IJCNN), 2025. <https://ieeexplore.ieee.org/document/11181772/>
- [7].T. Lin, "Structured retrieval-augmented generation for multi-entity question answering over heterogeneous sources," in Proc. 2025 Int. Joint Conf. Neural Networks (IJCNN), 2025. <https://ieeexplore.ieee.org/document/11107459/>
- [8].Z. Hammane, F.-E. Ben-Bouazza, and A. Fennan, "SelfRewardRAG: Enhancing medical reasoning with retrieval-augmented generation and self-evaluation in large language models," in Proc. 2024 Int. Conf. Inventive Computation Technologies (ICICT), 2024. <https://ieeexplore.ieee.org/document/10620139/>
- [9].J. Yao, Y. Guo, and Q. Yu, "Adapting large language models for healthcare with an enhanced retrieval-augmented generation framework," in Proc. 2024 IEEE Int. Conf. Bioinformatics and Biomedicine (BIBM), 2024. <https://ieeexplore.ieee.org/document/11019949/>

- [10]. S. M. H. Hashemi, H. B. Kashani, and S. Fatemi, "MediMind: A domain-adaptive medical LLM leveraging fine-tuning and retrieval-augmented generation," in Proc. 2024 IEEE Int. Conf. Systems, Man, and Cybernetics (SMC), 2024. <https://ieeexplore.ieee.org/document/11297589/>
- [11]. L. Stuhlmann, M. A. Saxer, and J. Fürst, "Efficient and reproducible biomedical question answering using retrieval augmented generation," in Proc. 2024 IEEE Int. Conf. Bioinformatics and Biomedicine (BIBM), 2024. <https://ieeexplore.ieee.org/document/11081505/>
- [12]. S. Guglani, A. K. Jakhar, A. Dasgupta, and S. Roy, "A RAG-based medical assistant especially for infectious diseases," in Proc. 2024 Int. Conf. Inventive Computation Technologies (ICICT), 2024. <https://ieeexplore.ieee.org/document/10620221/>
- [13]. R. Ban, "RARoK: Retrieval-augmented reasoning on knowledge for medical question answering," in Proc. 2024 IEEE Int. Conf. Bioinformatics and Biomedicine (BIBM), 2024. <https://ieeexplore.ieee.org/document/10984872/>
- [14]. B. Saha, U. Saha, and M. Z. Malik, "QuIM-RAG: Advancing retrieval-augmented generation with inverted question matching for enhanced QA performance," IEEE Access, vol. 12, 2024, doi: 10.1109/ACCESS.2024.3513155. <https://ieeexplore.ieee.org/document/10781379/>
- [15]. K. Sawarkar, A. Mangal, and S. R. Solanki, "Blended RAG: Improving RAG accuracy with semantic search and hybrid query-based retrievers," in Proc. 2024 IEEE Int. Conf. Multimedia Information Processing and Retrieval (MIPR), 2024. <https://ieeexplore.ieee.org/document/10707868/>
- [16]. H. M. A. Zeeshan, M. Faizan, U. Zia, and A. Gohar, "RAG powered LLMs for QA: Evolution, challenges, applications, and future directions," in Proc. 2024 IEEE Int. Computer Software and Applications Conf. (COMPSAC), 2024. <https://ieeexplore.ieee.org/document/11034531/>
- [17]. X. Xu, D. Zhang, W. Zhang, Q. Lu, and L. Zhu, "Design process for retrieval augmented generation systems," in Proc. 2024 IEEE Int. Conf. Software and System Processes (ICSSP), 2024. <https://ieeexplore.ieee.org/document/11015035/>
- [18]. T. Hecking, T. Sommer, and M. Felderer, "An architecture and protocol for decentralized retrieval augmented generation," in Proc. 2024 IEEE Int. Computer Software and Applications Conf. (COMPSAC), 2024. <https://ieeexplore.ieee.org/document/11014986/>
- [19]. M. A. Shavaki, P. Omrani, R. Toosi, and M. A. Akhaee, "Knowledge graph based retrieval-augmented generation for multi-hop question answering enhancement," in Proc. 2024 IEEE Int. Conf. Parallel and Distributed Systems (ICPADS), 2024. <https://ieeexplore.ieee.org/document/10892619/>

- [20]. S. Mankari and A. Sanghavi, "Enhancing vector based retrieval augmented generation with contextual knowledge graph construction," in Proc. 2024 IEEE Int. Conf. Big Data (BigData), 2024. <https://ieeexplore.ieee.org/document/10842699/>
- [21]. C. Lee, D. Prahlad, D. Kim, and H. Kim, "Work-in-progress: On-device retrieval augmented generation with knowledge graphs for personalized large language models," in Proc. 2024 Int. Conf. Compilers, Architecture, and Synthesis for Embedded Systems (CASES), 2024. <https://ieeexplore.ieee.org/document/10740767/>
- [22]. R. Koç, M. K. Gürkan, and F. T. Yarman Vural, "ReRag: A new architecture for reducing the hallucination by retrieval-augmented generation," in Proc. 2024 IEEE 32nd Signal Processing and Communications Applications Conf. (SIU), 2024. <https://ieeexplore.ieee.org/document/10773428/>
- [23]. S. AboulEla, P. Zabihitari, N. Ibrahim, M. Afshar, and R. Kashef, "Exploring RAG solutions to reduce hallucinations in LLMs," in Proc. 2024 IEEE Int. Computer Software and Applications Conf. (COMPSAC), 2024. <https://ieeexplore.ieee.org/document/11014810/>
- [24]. B. Malin, T. Kalganova, and N. Boulgouris, "A review of faithfulness metrics for hallucination assessment in large language models," IEEE Access, vol. 12, 2024, doi: 10.1109/ACCESS.2024.3488601. <https://ieeexplore.ieee.org/document/11032180/>
- [25]. S. Knollmeyer, M. U. Akmal, L. Koval, S. Asif, S. G. Mathias, and D. Großmann, "Document knowledge graph to enhance question answering with retrieval augmented generation," in Proc. 2024 IEEE 29th Int. Conf. Emerging Technologies and Factory Automation (ETFA), 2024. <https://ieeexplore.ieee.org/document/10711054/>
- [26]. P. Joshi, A. Gupta, P. Kumar, and M. Sisodia, "Robust multi model RAG pipeline for documents containing text, table & images," in Proc. 2024 IEEE 44th Int. Conf. Distributed Computing Systems (ICDCS), 2024. <https://ieeexplore.ieee.org/document/10574972/>
- [27]. R. Zhang, H. Du, Y. Liu, D. Niyato, J. Kang, and S. Sun, "Interactive AI with retrieval-augmented generation for next generation networking," IEEE Network, vol. 38, no. 4, pp. 144–151, 2024, doi: 10.1109/MNET.2024.10531073. <https://ieeexplore.ieee.org/document/10531073/>
- [28]. Sanjay Chakraborty, Hrithik Paul, Sayani Ghatak, Saroj Kumar Pandey, Ankit Kumar, Kamred Udham Singh, Mohd Asif Shah, "An AI-Based Medical Chatbot Model for Infectious Disease Prediction," IEEE Access, vol. 10, 2023. <https://ieeexplore.ieee.org/document/9970731>
- [29]. S. Sutton, A. Pincock, D. Baumgart, D. Bernstein, R. Bhatt, and K. Buchan, "A Systematic Review of Technologies and Standards Used in the Development of Rule-Based Clinical Decision Support Systems," Health Technol., vol. 12, no. 4, 2022. <https://link.springer.com/article/10.1007/s12553-022-00672-9>

- [30]. Mahmoud Nasr, Md. Milon Islam, Shady Shehata, Fakhri Karray and Yuri Quintana, "Smart Healthcare in the Age of AI: Recent Advances, Challenges, and Future Prospects," *IEEE Access*, vol. 9, 2021. <https://ieeexplore.ieee.org/document/9565155>

Base Paper:

From References the mainly referred paper: [26] Joshi, P., Gupta, A., Kumar, P. and Sisodia, M., 2024, June. Robust multi model rag pipeline for documents containing text, table & images. In *2024 3rd International Conference on Applied Artificial Intelligence and Computing (ICAAIC)* (pp. 993-999). IEEE.

APPENDIX

1. Publications



Fig A.1: Research Paper Publication

2. GitHub Link :

- GitHub: <https://github.com/aslamaslam1842004/ISE-28-rural-med-nexus-an-offline-intelligent-medical-knowledge-system>

3. Live Demo:

- Link: [Demo Link](#)

4. Screenshots




Fig A.2: Medical Data Ingestion Interface




Fig A.3: Select file for uploading medical data




Fig A.4: Uploading Successful File creates an Ingestion Trigger




Fig A.5: Execution of Ingestion Pipeline and the Processing Output


```

Problems  Output  Debug Console  Terminal  Ports
(venv) PS D:\8TH SEM\testing-rag\rag_enhanced\rag> python -m streamlit run ui/ingestion_app.py
>>
Loading documents: 100%| 27/27 [46:41<00:00, 103.75s/it]
2026-04-15 01:12:19.870 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:32 - Loaded 25 doc
ument(s)
2026-04-15 01:12:19.871 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:36 - Step 2: Loadi
ng images
2026-04-15 01:12:19.871 | INFO | src.ingestion.image_loader:load_images:42 - Found 2 image(s)
2026-04-15 01:12:22.239 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:38 - Loaded 2 imag
2026-04-15 01:12:22.239 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:38 - Loaded 2 imag
e(s)
2026-04-15 01:12:22.239 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:42 - Step 3: Loadi
2026-04-15 01:12:22.239 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:42 - Step 3: Loadi
ng audio
2026-04-15 01:12:22.239 | WARNING | src.ingestion.audio_loader:load_audio:44 - No audio files found.
Skipping audio ingestion.
2026-04-15 01:12:22.239 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:44 - Loaded 0 audi
o file(s)
2026-04-15 01:12:22.241 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:53 - Step 4: Chunk
ing text data
2026-04-15 01:12:22.241 | INFO | src.preprocessing.chunker:chunk_documents:63 - Chunking 35 docume
nt(s)
2026-04-15 01:12:24.054 | INFO | src.preprocessing.chunker:chunk_documents:88 - Generated 2366 chu
nks
2026-04-15 01:12:24.054 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:55 - Generated 236
6 chunks
2026-04-15 01:12:24.054 | INFO | src.ingestion.ingestion:run_ingestion_pipeline:63 - Step 5: Build
ing FAISS index
2026-04-15 01:12:24.054 | INFO | src.embeddings.vector_store:build_faiss_index:25 - Loading embedd

```

Fig A.6: Data Loading and Preprocessing Execution




Fig A.7: Completion of ingestion process and update of knowledge base

```
PS D:\8TH SEM\testing-rag\rag_enhanced\rag> .\venv\Scripts\activate
(venv) PS D:\8TH SEM\testing-rag\rag_enhanced\rag> python -m streamlit run ui/query_app.py
>>

You can now view your Streamlit app in your browser.

Local URL: http://localhost:8502
Network URL: http://192.168.1.8:8502
```

Fig A.8: Start Query Interface with Streamlit




Fig A .9: Selection interface of medical query




Fig A.10: Screenshot of Query Response and Source Retrieval Interface